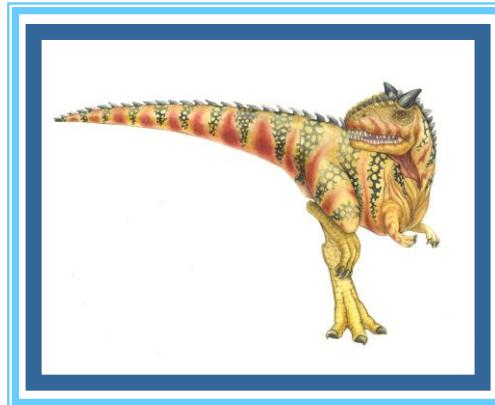


Chapter 4: Threads





Chapter 4: Threads

- Overview
- Multicore Programming
- Multithreading Models
- Thread Libraries
- Implicit Threading
- Threading Issues





Motivation

- Most modern applications are multithreaded
- Threads run within application
- ***Process creation is heavy-weight while thread creation is light-weight***
- Can simplify code, increase efficiency
- **Kernels are generally multithreaded**






What is?

- A thread is a basic unit of CPU utilization; **a fundamental unit of CPU utilization** that forms the basis of multithreaded computer systems
- It comprises
 - a thread ID,
 - a program counter,
 - a register set,
 - a stack.
- It shares with other threads belonging to the same process
 - its **code section**,
 - **data section**,
 - and other operating-system resources, such as **open files** and **signals**.






Single and Multithreaded Processes



single-threaded process



multithreaded process





Benefits

- **Responsiveness** – may allow continued execution if part of process is blocked, *especially important for user interfaces*
- **Resource Sharing** – threads **share** resources of process, easier than *shared memory* or *message passing*
 - **SM** and **MP** techniques must be explicitly arranged by the programmer
 - Threads **share** the **memory** and the **resources** of the process **by default**
- **Economy** – cheaper than process creation, **thread switching** lower overhead than **context switching**
 - In Solaris, creating a process is about **thirty times slower** than is creating a thread, and context switching is about **five times slower**.
- **Scalability** – process can take advantage of multiprocessor architectures





4.2 Multicore Programming

- Each core appears as a separate processor to the operating system
- **Multicore** or **multiprocessor** systems putting pressure on programmers, **challenges** include:
 - **Identifying tasks**
 - ▶ This involves examining applications to find areas that can be divided into separate, concurrent tasks
 - **Data splitting**
 - ▶ Data accessed and manipulated by the tasks must be divided to run on separate cores
 - **Balance**
 - ▶ Tasks should perform **equal work of equal value**.
 - **Data dependency**
 - ▶ The data accessed by the tasks must be examined for dependencies between two or more tasks. Execution of the tasks is **synchronized** to accommodate the data dependency.





4.2 Multicore Programming


- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
 - The execution of the threads will be **interleaved over time**
 - Single processor / core, scheduler providing concurrency
 - In non-SMP systems (Single core systems) CPU schedulers were designed to provide the illusion of parallelism by rapidly switching between processes






Concurrency vs. Parallelism

- Concurrent execution on single-core system:




- Parallelism on a multi-core system:





Multicore Programming (Cont.)

- As # of threads grows, so does architectural support for threading





Amdahl's Law

- Identifies performance gains from **adding additional cores** to an application that has both serial and parallel components
- S is serial portion
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$

Serial portion of an application has disproportionate effect on performance gained by adding additional cores





Multicore Programming (Cont.)

□ Types of parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - ▶ For example, summing the contents of an array of size N . On a single-core system, one thread would simply sum the elements $[0] \dots [N - 1]$.
 - ▶ On a dual-core system, however, thread *A*, running on core 0, could sum the elements $[0] \dots [N/2 - 1]$ while thread *B*, running on core 1, could sum the elements $[N/2] \dots [N - 1]$.
- **Task parallelism** – distributing threads across cores, each thread performing unique operation
 - ▶ For example, two threads, each performing a unique statistical operation on the array of elements. The threads again are operating in parallel on separate computing cores, but each is performing a **unique operation**.





4.3 Multithreading Models

User Threads and Kernel Threads

- **User threads** - management done by user-level threads library
 - Three primary thread libraries:
 - ▶ POSIX **Pthreads**
 - ▶ Windows threads
 - ▶ Java threads
- **Kernel threads** - Supported by the Kernel and managed by the OS
 - Examples – virtually all general purpose operating systems, including:
 - ▶ Windows
 - ▶ Solaris
 - ▶ Linux
 - ▶ Tru64 UNIX
 - ▶ Mac OS X





4.3 Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many





Many-to-One


- Many user-level threads mapped to single kernel thread
- The entire process will **block** if a thread makes a **blocking system call**.
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
 - **Solaris Green Threads**
 - **GNU Portable Threads**





One-to-One


- Each user-level thread maps to kernel thread
- **Creating a user-level thread creates a kernel thread**
- **More concurrency than many-to-one**
 - It also allows multiple threads to run in parallel on multiprocessors
- Number of threads per process sometimes restricted due to overhead
- Examples
 - Windows
 - Linux
 - Solaris 9 and later





Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- When a thread performs a blocking system call, the **kernel can schedule another thread for execution**
- Solaris prior to version 9
- Windows with the *ThreadFiber* package





Thread Libraries

- For **POSIX** and Windows threading, any data declared **globally**—that is, declared outside of any function—are **shared** among all threads belonging to the same process.
- **Java** has **no notion of global data**, access to shared data must be **explicitly arranged between threads**
- Data declared local to a function are typically stored on the **stack**
- Since each thread has its own stack, each thread has its own copy of local data.





Asynchronous vs synchronous threading

□ Asynchronous threading

- Once the parent creates a child thread, the parent resumes its execution, so that the parent and child **execute concurrently**
- Each thread runs independently of every other thread
 - ▶ E.g. Multithreaded web server

□ Synchronous threading

- Parent thread creates one or more children and then must wait for all of its children to terminate before it resumes
- Known as **fork-join strategy**
- **Threads can run concurrently but parent cannot continue until this work has been completed**
- Synchronous threading involves **significant data sharing** among threads





PThreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
 - *It is a specification*, not an *implementation*
- API specifies **behavior** of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)





Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr,"usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr,"%d must be >= 0\n",atoi(argv[1]));
        return -1;
    }
}
```





PThreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

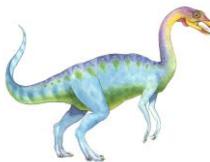
printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```





Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```





Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* the thread runs in this separate function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 0; i <= Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    if (argc != 2) {
        fprintf(stderr,"An integer parameter is required\n");
        return -1;
    }
    Param = atoi(argv[1]);
    if (Param < 0) {
        fprintf(stderr,"An integer >= 0 is required\n");
        return -1;
    }
}
```





Windows Multithreaded C Program (Cont.)

```
/* create the thread */
ThreadHandle = CreateThread(
    NULL, /* default security attributes */
    0, /* default stack size */
    Summation, /* thread function */
    &Param, /* parameter to thread function */
    0, /* default creation flags */
    &ThreadId); /* returns the thread identifier */

if (ThreadHandle != NULL) {
    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n", Sum);
}

}
```





Java Threads

- Java threads are managed by the JVM
- Typically implemented using the threads model provided by underlying OS
- Two way of implementing Threads in Java
 - Extending Thread class
 - ▶ Create a new class that is derived from the **Thread** class and to override its **run()** method
 - Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```

- ▶ When a class implements **Runnable**, it must define a **run()** method. The code implementing the **run()** method is what runs as a separate thread





Java Multithreaded Program

```
class Sum
{
    private int sum;

    public int getSum() {
        return sum;
    }

    public void setSum(int sum) {
        this.sum = sum;
    }
}

class Summation implements Runnable
{
    private int upper;
    private Sum sumValue;

    public Summation(int upper, Sum sumValue) {
        this.upper = upper;
        this.sumValue = sumValue;
    }

    public void run() {
        int sum = 0;
        for (int i = 0; i <= upper; i++)
            sum += i;
        sumValue.setSum(sum);
    }
}
```





Java Multithreaded Program (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        if (args.length > 0) {
            if (Integer.parseInt(args[0]) < 0)
                System.err.println(args[0] + " must be >= 0.");
            else {
                Sum sumObject = new Sum();
                int upper = Integer.parseInt(args[0]);
                Thread thrd = new Thread(new Summation(upper, sumObject));
                thrd.start(); // Call to start() method
                try {
                    thrd.join();
                    System.out.println
                        ("The sum of "+upper+" is "+sumObject.getSum());
                } catch (InterruptedException ie) { }
            }
        }
        else
            System.err.println("Usage: Summation <integer value>"); }
    }
```

Creating a **Thread** object does not specifically create the new thread; rather, the **start()** method creates the new thread

It calls the **run()** method, making the thread eligible to be run by the JVM.
(Note again that we never call the **run()** method directly.





Implicit Threading

- Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- Creation and management of threads done by compilers and run-time libraries rather than programmers
- Three methods
 - Thread Pools
 - OpenMP
 - Grand Central Dispatch
- Other methods include Intel's Threading Building Blocks (TBB), `java.util.concurrent` package





OpenMP

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in **shared-memory environments**
- Identifies **parallel regions** – blocks of code that can run in parallel

```
#pragma omp parallel
```

Create as many threads as there are **cores**

```
#pragma omp parallel for
for(i=0;i<N;i++) {
    c[i] = a[i] + b[i];
}
```

Run for loop in parallel

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */

    return 0;
}
```





Threading Issues

- **Issue:** The semantics of the **fork()** and **exec()** system calls **change** in a multithreaded program.
 - If one thread in a program calls **fork()**, does the new process duplicate all threads, or is the new process single-threaded?
- Signal handling
 - A **signal** is used in UNIX systems to notify a process that a particular **event has occurred**.
 - Synchronous and asynchronous
- Thread cancellation of target thread
 - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations





Thread Cancellation

- Terminating a thread before it has finished
- Thread to be canceled is **target thread**
- Two general approaches:
 - **Asynchronous cancellation** terminates the target thread immediately
 - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- Pthread code to create and cancel a thread:

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
.  
.  
.  
  
/* cancel the thread */  
pthread_cancel(tid);
```

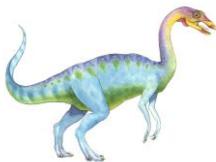




Thread-Local Storage

- **Thread-local storage (TLS)** allows each thread to have its own copy of data
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- **Different from local variables**
 - Local variables visible only during single function invocation
 - TLS visible across function invocations
- Similar to `static` data
 - **TLS is unique to each thread**
- The functions `pthread_key_create` and `pthread_key_delete` are used respectively to create and delete a key for thread-specific data.
- The type of the key is explicitly left opaque and is referred to as `pthread_key_t`.

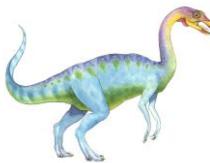




News in C11

- Multi-threading support (_Thread_local storage-class specifier, <threads.h> header including thread creation/management functions, mutex, condition variable and thread-specific storage functionality)





News in C++11

- `#include <threads.h>`
- `thread_local int foo = 0;`
- C++11 introduces the `thread_local` keyword



End of Chapter 4

