
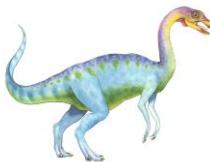


# Chapter 5: Process Synchronization





# Chapter 5: Process Synchronization

- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples
- Alternative Approaches

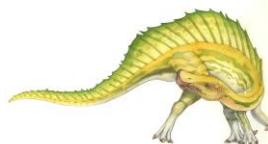




# Objectives

---

- To present the concept of process synchronization.
- To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
- To present both software and hardware solutions of the critical-section problem
- To examine several classical process-synchronization problems
- To explore several tools that are used to solve process synchronization problems



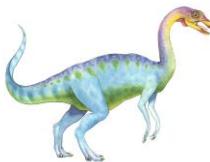


# Background

---

- Processes can execute concurrently
  - May be interrupted at any time, partially completing execution
- **Concurrent access** to **shared** data may result in **data inconsistency**
- **Illustration of the problem:**  
Suppose that we wanted to provide a solution to the consumer-producer problem that fills ***all*** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





# Producer

---

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

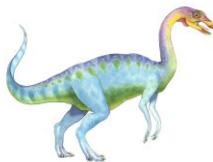




# Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```





# Race Condition

- `counter++` could be implemented as

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

- `counter--` could be implemented as

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

- Consider this execution interleaving with “count = 5” initially:

S0: producer execute <code>register1 = counter</code>	{register1 = 5}
S1: producer execute <code>register1 = register1 + 1</code>	{register1 = 6}
S2: consumer execute <code>register2 = counter</code>	{register2 = 5}
S3: consumer execute <code>register2 = register2 - 1</code>	{register2 = 4}
S4: producer execute <code>counter = register1</code>	{counter = 6 }
S5: consumer execute <code>counter = register2</code>	{counter = 4}

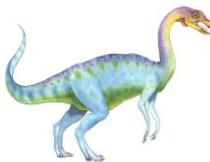




# Critical Section Problem

- Consider system of  $n$  processes  $\{p_0, p_1, \dots p_{n-1}\}$
- Each process has **critical section** segment of code
  - Process may be changing common variables, updating table, writing file, etc
  - **When one process in critical section, no other may be in its critical section**
- **Critical section problem** is to design a **protocol** that the processes can use to cooperate.
- Each process must ask **permission** to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





# Critical Section

---

- General structure of process  $P_i$

```
do {
```

```
    entry section
```

```
    critical section
```

```
    exit section
```

```
    remainder section
```

```
} while (true);
```






# Algorithm for Process $P_i$

```
do {  
  
    while (turn == j);  
  
        critical section  
  
    turn = j;  
  
        remainder section  
  
} while (true);
```





# Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then **no other processes** can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that **wish to enter** their critical section, then the selection of the processes that will enter the critical section next **cannot be postponed indefinitely**
3. **Bounded Waiting** - **A bound must exist** on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - Assume that each process executes at a nonzero speed
  - No assumption concerning **relative speed** of the  $n$  processes



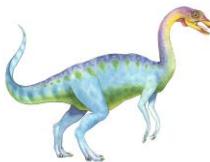


# Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive kernels**— allows preemption of process when running in kernel mode
  - ▶ **More responsive** but it **must be carefully designed**
- **Non-preemptive kernels** – runs until exits kernel mode, blocks, or voluntarily yields CPU
  - ▶ **Essentially free of race conditions in kernel mode**





# Peterson's Solution

---

- Good algorithmic description of solving the problem
- Two process solution
- **Assume that the `load` and `store` machine-language instructions are atomic**; that is, **cannot be interrupted**
- The two processes share two variables:
  - `int turn;`
  - `Boolean flag[2]`
- The variable `turn` indicates whose turn it is to enter the critical section
- The `flag` array is used to indicate if a process is ready to enter the critical section. `flag[i] = true` implies that process  $P_i$  is ready!





# Algorithm for Process $P_i$

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = false;  
        remainder section  
} while (true);
```





## Peterson's Solution (Cont.)

- Provable that the three CS requirement are met:
  1. Mutual exclusion is preserved
  2. Progress requirement is satisfied
  3. Bounded-waiting requirement is met






# Synchronization Hardware

---

- Many systems provide hardware support for implementing the critical section code.
- All solutions below based on idea of **locking**
  - Protecting critical regions via locks
- Uniprocessors – could disable interrupts
  - Currently running code would execute without preemption
  - Generally too inefficient on multiprocessor systems
    - ▶ Operating systems using this not broadly scalable
- Modern machines provide special atomic hardware instructions
  - ▶ **Atomic** = non-interruptible
    - Either test memory word and set value
    - Or swap contents of two memory words

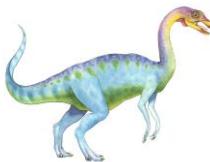




# Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```





# test\_and\_set Instruction

---

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to “TRUE”.





# Solution using test\_and\_set()

---

- Shared Boolean variable lock, initialized to FALSE
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */  
} while (true);
```





# Mutex Locks

---

- ❑ Previous solutions are **complicated** and generally **inaccessible** to application programmers
- ❑ OS designers build **software tools** to solve critical section problem
- ❑ Simplest is mutex lock
- ❑ Protect a critical section by first **acquire()** a lock then **release()** the lock
  - ❑ Boolean variable indicating if lock is available or not
- ❑ Calls to **acquire()** and **release()** must be atomic
  - ❑ Usually implemented via hardware atomic instructions
- ❑ But this solution requires **busy waiting**
  - ❑ This lock therefore called a **spinlock**





# acquire() and release()


```
□ acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}  
□ release() {  
    available = true;  
}  
□ do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```





# Semaphore

- Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- Semaphore **S** – **integer variable**
- Can only be accessed via two indivisible (atomic) operations
  - **wait()** and **signal()**
    - ▶ Originally called **P()** and **V()**



```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```





# Semaphore Usage

- **Counting semaphore** – integer value can range over an unrestricted domain
- **Binary semaphore** – integer value can range only between 0 and 1
  - Same as a **mutex lock**
- Can solve various synchronization problems
- Consider  $P_1$  and  $P_2$  that require  $S_1$  to happen before  $S_2$   
Create a semaphore “synch” **initialized to 0**

**P1:**

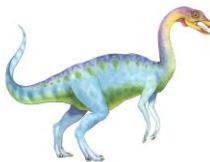
```
s1;  
    signal(synch);
```

**P2:**

```
wait(synch);  
s2;
```

- Can implement a counting semaphore  $S$  as a binary semaphore






# Deadlock and Starvation

- **Deadlock** – two or more processes are **waiting indefinitely** for an event that can be caused by only one of the waiting processes
- Let  $S$  and  $Q$  be two semaphores initialized to 1

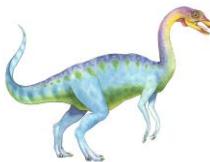
```
P0
    wait(S) ;
    wait(Q) ;
    ...
    signal(S) ;
    signal(Q) ;
```

```
P1
    wait(Q) ;
    wait(S) ;
    ...
    signal(Q) ;
    signal(S) ;
```



- **Starvation – indefinite blocking**
  - A process may never be removed from the semaphore queue in which it is suspended
- **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
  - Solved via **priority-inheritance protocol**





# Semaphore definition

```
typedef struct {
    int value;
    struct process *list;
} semaphore;

wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

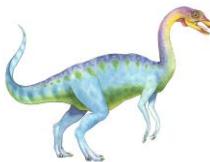
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```



# Classical Problems of Synchronization

- Classical problems used to test newly-proposed synchronization schemes
  - Bounded-Buffer Problem
  - Readers and Writers Problem
  - Dining-Philosophers Problem





# Bounded-Buffer Problem

---

- $n$  buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
- Semaphore **full** initialized to the value 0
- Semaphore **empty** initialized to the value n



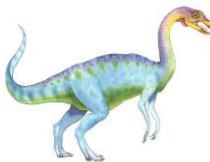


# Bounded Buffer Problem (Cont.)

- The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```



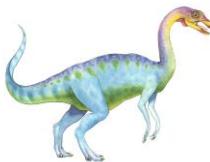


# Bounded Buffer Problem (Cont.)

- The structure of the consumer process

```
Do {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to  
next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```






# Readers-Writers Problem

---

- A data set is shared among a number of concurrent processes
  - Readers – **only read** the data set; they do **not** perform any updates
  - Writers – can **both read and write**
- **Problem – allow multiple readers to read at the same time**
  - Only one single writer can access the shared data at the same time
- Several variations of how readers and writers are considered – all involve some form of priorities
- Shared Data
  - Data set
  - Semaphore **rw\_mutex** initialized to 1
  - Semaphore **mutex** initialized to 1
  - Integer **read\_count** initialized to 0





# Readers-Writers Problem (Cont.)

- The structure of a writer process

```
do  {  
    wait(rw_mutex);  
  
    ...  
    /* writing is performed */  
  
    ...  
    signal(rw_mutex);  
} while (true);
```





## □ Writers

- Only one writer writes at a time
- While writing, reading is not possible

## □ Readers

- Multiple readers can read concurrently
- While reading, writing is not possible

Writer

```
Wait(rw_mutex)  
Write()  
Signal(rw_mutex)
```

Reader

```
Wait(mutex)  
read_count++;  
If(read_count==1)  
Wait(rw_mutex)  
Signal(mutex)
```

Read()

```
Wait(mutex)  
read_count--;  
If (read_count==0)  
Signal(rw_mutex)  
Signal(mutex)
```





# Readers Writers Problem

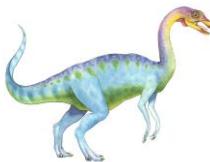
## Writers

Only one writer writes at a time  
While writing, reading is not possible

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible





# Readers Writers Problem

mutex=1

## Writers

Only one writer writes at a time  
While writing, reading is not possible

```
wait(mutex);  
write();  
signal(mutex);
```


## Readers

Multiple readers can read concurrently  
While reading, writing is not possible

```
wait(mutex);  
read();  
signal(mutex);
```

This solution allows only one reader or writer to access shared variable





# Readers Writers Problem

## Writers

Only one writer writes at a time  
While writing, reading is not possible

```
mutex=1
```

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible

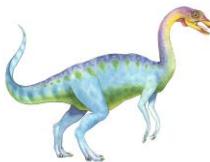
```
wait(mutex);  
write();  
signal(mutex);
```

```
wait(mutex);  
read();  
signal(mutex);
```

```
wait(mutex);  
read();  
signal(mutex);
```

Multiple readers cannot read





# Readers Writers Problem

## Writers

Only one writer writes at a time  
While writing, reading is not possible

rwmutex=1

```
wait(rwmutex);  
write();  
signal(rwmutex);
```

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible

```
readers++;  
If (readers==1)  
    wait(rwmutex);
```

```
read();
```

```
readers--;  
If(readers==0)  
    signal(rwmutex);
```

Multiple readers can read now  
However, accessing the «readers» variable is now creating a race condition



# Readers Writers Problem

## Writers

Only one writer writes at a time  
While writing, reading is not possible

rwmutex=1  
mutex=1

```
wait(rwmutex);  
write();  
signal(rwmutex);
```

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible


```
wait(mutex);  
readers++;  
If (readers==1)  
    wait(rwmutex);  
signal(mutex);  
read();
```

## Readers

Multiple readers can read concurrently  
While reading, writing is not possible

```
wait(mutex);  
readers--;  
If(readers==0)  
    signal(rwmutex);  
signal(mutex);
```

Problem solved



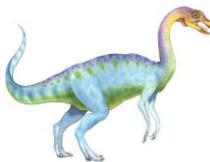


# Readers-Writers Problem (Cont.)

- The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
    ...  
    /* reading is performed */  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```



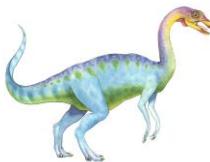


# Dining-Philosophers Problem



- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (**one at a time**) to eat from bowl
  - Need both to eat, then release both when done
- In the case of 5 philosophers
  - Shared data
    - ▶ Bowl of rice (data set)
    - ▶ Semaphore **chopstick [5]** initialized to 1





# Dining-Philosophers Problem Algorithm

- The structure of Philosopher  $i$ :

```
do {  
    wait (chopstick[i] );  
    wait (chopStick[ (i + 1) % 5] );  
  
        // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5] );  
  
        // think  
  
} while (TRUE);
```

- What is the problem with this algorithm?

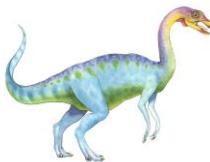




# Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock handling
  - Allow at most 4 philosophers to be sitting simultaneously at the table.
  - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
  - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.





# Problems with Semaphores

- Incorrect use of semaphore operations:
  - signal (mutex) .... wait (mutex) \* unless you want sequential processing
  - wait (mutex) ... wait (mutex)
  - Omitting of wait (mutex) or signal (mutex) (or both)
- Deadlock and starvation are possible.





# Monitors

---


- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- *Abstract data type*, internal variables only accessible by code within the procedure
- Only one process may be active within the monitor at a time
- But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }


    procedure Pn (...) { ..... }

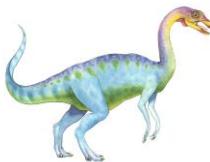
    Initialization code (...) { ... }
}
```





# Schematic view of a Monitor







# Condition Variables

- **condition x, y;**
- Two operations are allowed on a condition variable:
  - **x.wait()** – a process that invokes the operation is suspended until **x.signal()**
  - **x.signal()** – resumes one of processes (if any) that invoked **x.wait()**
    - ▶ If no **x.wait()** on the variable, then it has no effect on the variable





# Monitor with Condition Variables






# Condition Variables

---

- Condition variables **are not semaphores**. They are different even though they look similar.
- A condition variable does not count: have no associated integer.
- A signal on a condition variable  $x$  is **lost** (not saved for future use) if there is no process waiting (blocked) on the condition variable  $x$ .
- The **wait()** operation on a condition variable  $x$  will **always cause the caller of wait to block**.
- The **signal()** operation on a condition variable will wake up a sleeping process on the condition variable, **if any**. It has no effect if there is nobody sleeping.






# Condition variables: example

- Assume we have a resource to be accessed by many processes. Assume we have 5 instances of the resource. 5 processes can use the resource simultaneously.
  
- We want to implement a monitor that will implement two functions: acquire() and release() that can be called by a process before and after using a resource.





# Condition variables: example

**monitor** Allocate

{

    int count = 5; // we initialize count to 5.

    condition c;

    void acquire () {

        if (count == 0)

            c.wait();

        count--;

    }

    void release () {

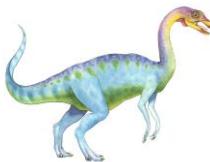
        count++;

        c.signal();

    }

}





# A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```





# Dining Philosophers

```
monitor DiningPhilosophers
{
    enum {THINKING, HUNGRY, EATING} state[5];
    condition self[5];
    void pickup(int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING)
            self[i].wait();
    }
    void putdown(int i) {
        state[i] = THINKING;
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
    void test(int i) {
        if ((state[(i + 4) % 5] != EATING) &&
            (state[i] == HUNGRY) &&
            (state[(i + 1) % 5] != EATING)) {
            state[i] = EATING;
            self[i].signal();
        }
    }
    initialization code() {
        for (int i = 0; i < 5; i++)
            state[i] = THINKING;
    }
}
```

DiningPhilosophers.pickup(i);  
...  
eat  
...  
DiningPhilosophers.putdown(i);



# End of Chapter 5

