


Chapter 5: Process Synchronization





Chapter 5: Process Synchronization

- n Background
- n The Critical-Section Problem
- n Peterson's Solution
- n Synchronization Hardware
- n Mutex Locks
- n Semaphores
- n Classic Problems of Synchronization
- n Monitors
- n Synchronization Examples
- n Alternative Approaches

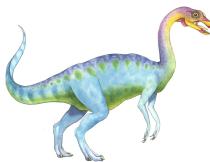




Objectives

- n To present the concept of process synchronization.
- n To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
- n To present both software and hardware solutions of the critical-section problem
- n To examine several classical process-synchronization problems
- n To explore several tools that are used to solve process synchronization problems





Background

- n Processes can execute concurrently
 - | May be interrupted at any time, partially completing execution
- n Concurrent access to shared data may result in data inconsistency
- n Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes
- n Illustration of the problem:
Suppose that we wanted to provide a solution to the consumer-producer problem that fills ***all*** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





Producer

```
while (true) {  
    /* produce an item in next_produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```





Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next_consumed */  
}
```





Race Condition

- n `counter++` could be implemented as

```
register1 = counter  
register1 = register1 + 1  
counter = register1
```

- n `counter--` could be implemented as

```
register2 = counter  
register2 = register2 - 1  
counter = register2
```

- n Consider this execution interleaving with “count = 5” initially:

- n S0: producer execute `register1 = counter` {register1 = 5}
S1: producer execute `register1 = register1 + 1` {register1 = 6}
S2: consumer execute `register2 = counter` {register2 = 5}
S3: consumer execute `register2 = register2 - 1` {register2 = 4}
S4: producer execute `counter = register1` {counter = 6 }
S5: consumer execute `counter = register2` {counter = 4}

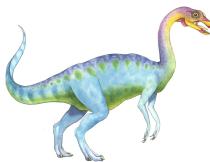




Critical Section Problem

- n Consider system of n processes $\{p_0, p_1, \dots, p_{n-1}\}$
- n Each process has **critical section** segment of code
 - | Process may be changing common variables, updating table, writing file, etc
 - | When one process in critical section, no other may be in its critical section
- n **Critical section problem** is to design protocol to solve this
- n Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**





Critical Section

- n General structure of process P_i

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```





Algorithm for Process P_i

```
do {  
  
    while (turn == j);  
        critical section  
  
    turn = j;  
        remainder section  
  
} while (true);
```





Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process P_i is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
 - Assume that each process executes at a nonzero speed
 - No assumption concerning **relative speed** of the n processes





Critical-Section Handling in OS

Two approaches depending on if kernel is preemptive or non-preemptive

- **Preemptive** – allows preemption of process when running in kernel mode
- **Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU
 - Essentially free of race conditions in kernel mode





Peterson's Solution

- n Good algorithmic description of solving the problem
- n Two process solution
- n Assume that the **load** and **store** machine-language instructions are atomic; that is, cannot be interrupted
- n The two processes share two variables:
 - | **int turn;**
 - | **Boolean flag[2]**
- n The variable **turn** indicates whose turn it is to enter the critical section
- n The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process **P_i** is ready!

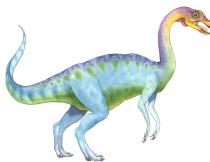




Algorithm for Process P_i

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = false;  
        remainder section  
} while (true);
```





Peterson's Solution (Cont.)

n Provable that the three CS requirement are met:

1. Mutual exclusion is preserved

P_i enters CS only if:

either `flag[j] = false` or `turn = i`

2. Progress requirement is satisfied
3. Bounded-waiting requirement is met





Synchronization Hardware

- n Many systems provide hardware support for implementing the critical section code.
- n All solutions below based on idea of **locking**
 - | Protecting critical regions via locks
- n Uniprocessors – could disable interrupts
 - | Currently running code would execute without preemption
 - | Generally too inefficient on multiprocessor systems
 - ▶ Operating systems using this not broadly scalable
- n Modern machines provide special atomic hardware instructions
 - ▶ **Atomic** = non-interruptible
 - | Either test memory word and set value
 - | Or swap contents of two memory words

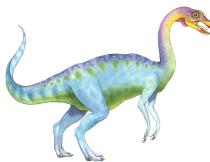




Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```





test_and_set Instruction

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```

1. Executed atomically
2. Returns the original value of passed parameter
3. Set the new value of passed parameter to “TRUE”.





Solution using test_and_set()

- n Shared Boolean variable lock, initialized to FALSE
- n Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
        /* critical section */  
    lock = false;  
        /* remainder section */  
} while (true);
```





compare_and_swap Instruction

Definition:

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
  
    if (*value == expected)  
        *value = new_value;  
  
    return temp;  
}
```

1. Executed atomically
2. Returns the original value of passed parameter “value”
3. Set the variable “value” the value of the passed parameter “new_value” but only if “value” ==“expected”. That is, the swap takes place only under this condition.





Solution using compare_and_swap

- n Shared integer “lock” initialized to 0;
- n Solution:

```
do {
    while (compare_and_swap(&lock, 0, 1) != 0)
        ; /* do nothing */
    /* critical section */
lock = 0;
    /* remainder section */
} while (true);
```





Bounded-waiting Mutual Exclusion with test_and_set

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /* remainder section */
} while (true);
```

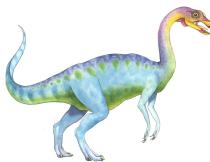




Mutex Locks

- n Previous solutions are complicated and generally inaccessible to application programmers
- n OS designers build software tools to solve critical section problem
- n Simplest is mutex lock
- n Protect a critical section by first **acquire()** a lock then **release()** the lock
 - | Boolean variable indicating if lock is available or not
- n Calls to **acquire()** and **release()** must be atomic
 - | Usually implemented via hardware atomic instructions
- n But this solution requires **busy waiting**
 - | This lock therefore called a **spinlock**






acquire() and release()

```
n   acquire() {
        while (!available)
n           ; /* busy wait */
        available = false;
    }
release() {
n     available = true;
}
do {
n     acquire lock
        critical section
    release lock
        remainder section
} while (true);
```





Semaphore

- n Synchronization tool that provides more sophisticated ways (than Mutex locks) for process to synchronize their activities.
- n Semaphore **S** – integer variable
- n Can only be accessed via two indivisible (atomic) operations
 - | **wait()** and **signal()**
 - ▶ Originally called **P()** and **V()**

- n Definition of the **wait()** operation

```
wait(S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}
```

- n Definition of the **signal()** operation

```
signal(S) {  
    S++;  
}
```





Semaphore Usage

- n **Counting semaphore** – integer value can range over an unrestricted domain
- n **Binary semaphore** – integer value can range only between 0 and 1
 - | Same as a **mutex lock**
- n Can solve various synchronization problems
- n Consider P_1 and P_2 that require S_1 to happen before S_2
Create a semaphore “**synch**” initialized to 0

P1:

```
S1;  
signal(synch);
```

P2:

```
wait(synch);  
S2;
```

- n Can implement a counting semaphore **S** as a binary semaphore





Semaphore Implementation

- n Must guarantee that no two processes can execute the **wait()** and **signal()** on the same semaphore at the same time
- n Thus, the implementation becomes the critical section problem where the **wait** and **signal** code are placed in the critical section
 - | Could now have **busy waiting** in critical section implementation
 - ▶ But implementation code is short
 - ▶ Little busy waiting if critical section rarely occupied
- n Note that applications may spend lots of time in critical sections and therefore this is not a good solution





Semaphore Implementation with no Busy waiting

- n With each semaphore there is an associated waiting queue
- n Each entry in a waiting queue has two data items:
 - | value (of type integer)
 - | pointer to next record in the list
- n Two operations:
 - | **block** – place the process invoking the operation on the appropriate waiting queue
 - | **wakeup** – remove one of processes in the waiting queue and place it in the ready queue
- n `typedef struct{
 int value;
 struct process *list;
} semaphore;`





Implementation with no Busy waiting (Cont.)

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}  
  
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```





Deadlock and Starvation

- n **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- n Let S and Q be two semaphores initialized to 1

P_0	P_1
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
...	...
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- n **Starvation – indefinite blocking**
 - | A process may never be removed from the semaphore queue in which it is suspended
- n **Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - | Solved via **priority-inheritance protocol**





Classical Problems of Synchronization

- n Classical problems used to test newly-proposed synchronization schemes
 - | Bounded-Buffer Problem
 - | Readers and Writers Problem
 - | Dining-Philosophers Problem





Bounded-Buffer Problem

- n n buffers, each can hold one item
- n Semaphore **mutex** initialized to the value 1
- n Semaphore **full** initialized to the value 0
- n Semaphore **empty** initialized to the value n





Bounded Buffer Problem (Cont.)

- n The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





Bounded Buffer Problem (Cont.)

- n The structure of the consumer process

```
Do {  
    wait(full);  
    wait(mutex);  
  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```

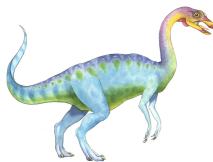




Readers-Writers Problem

- n A data set is shared among a number of concurrent processes
 - | Readers – only read the data set; they do *not* perform any updates
 - | Writers – can both read and write
- n Problem – allow multiple readers to read at the same time
 - | Only one single writer can access the shared data at the same time
- n Several variations of how readers and writers are considered – all involve some form of priorities
- n Shared Data
 - | Data set
 - | Semaphore **rw_mutex** initialized to 1
 - | Semaphore **mutex** initialized to 1
 - | Integer **read_count** initialized to 0





Readers-Writers Problem (Cont.)

- n The structure of a writer process

```
do {  
    wait(rw_mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw_mutex);  
} while (true);
```





Readers-Writers Problem (Cont.)

- n The structure of a reader process

```
do {  
    wait(mutex);  
    read_count++;  
    if (read_count == 1)  
        wait(rw_mutex);  
    signal(mutex);  
  
    ...  
    /* reading is performed */  
  
    ...  
    wait(mutex);  
    read_count--;  
    if (read_count == 0)  
        signal(rw_mutex);  
    signal(mutex);  
} while (true);
```






Readers-Writers Problem Variations

- n **First** variation – no reader kept waiting unless writer has permission to use shared object
- n **Second** variation – once writer is ready, it performs the write ASAP
- n Both may have starvation leading to even more variations
- n Problem is solved on some systems by kernel providing reader-writer locks



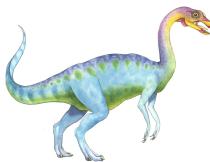


Dining-Philosophers Problem



- Philosophers spend their lives alternating thinking and eating
- Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl
 - | Need both to eat, then release both when done
- In the case of 5 philosophers
 - | Shared data
 - ▶ Bowl of rice (data set)
 - ▶ Semaphore **chopstick [5]** initialized to 1





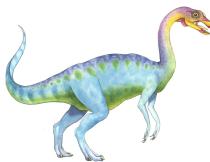
Dining-Philosophers Problem Algorithm

- n The structure of Philosopher *i*:

```
do {  
    wait (chopstick[i] ) ;  
    wait (chopStick[ (i + 1) % 5] ) ;  
  
        // eat  
  
    signal (chopstick[i] ) ;  
    signal (chopstick[ (i + 1) % 5] ) ;  
  
        // think  
  
} while (TRUE) ;
```

- n What is the problem with this algorithm?





Dining-Philosophers Problem Algorithm (Cont.)

n Deadlock handling

- | Allow at most 4 philosophers to be sitting simultaneously at the table.
- | Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
- | Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.






Problems with Semaphores

- n Incorrect use of semaphore operations:
 - | signal (mutex) wait (mutex)
 - | wait (mutex) ... wait (mutex)
 - | Omitting of wait (mutex) or signal (mutex) (or both)
- n Deadlock and starvation are possible.





Monitors

- n A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- n *Abstract data type*, internal variables only accessible by code within the procedure
- n Only one process may be active within the monitor at a time
- n But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }


    procedure Pn (...) { ..... }

    Initialization code (...) { ... }
}
```





Schematic view of a Monitor






Condition Variables

- n **condition x, y;**
- n Two operations are allowed on a condition variable:
 - | **x.wait()** – a process that invokes the operation is suspended until **x.signal()**
 - | **x.signal()** – resumes one of processes (if any) that invoked **x.wait()**
 - ▶ If no **x.wait()** on the variable, then it has no effect on the variable





Monitor with Condition Variables





Condition Variables Choices

- n If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
 - | Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- n Options include
 - | **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - | **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - | Both have pros and cons – language implementer can decide
 - | Monitors implemented in Concurrent Pascal compromise
 - ▶ P executing signal immediately leaves the monitor, Q is resumed
 - | Implemented in other languages including Mesa, C#, Java





Monitor Solution to Dining Philosophers

```
monitor DiningPhilosophers
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self[i].wait;
    }

    void putdown (int i) {
        state[i] = THINKING;
            // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```





Solution to Dining Philosophers (Cont.)

```
void test (int i) {  
    if ((state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}  
  
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}  
}
```





Solution to Dining Philosophers (Cont.)

- n Each philosopher i invokes the operations **pickup ()** and **putdown ()** in the following sequence:

DiningPhilosophers.pickup(i) ;

EAT

DiningPhilosophers.putdown(i) ;

- n No deadlock, but starvation is possible





Monitor Implementation Using Semaphores

- n Variables

```
semaphore mutex; // (initially = 1)
semaphore next; // (initially = 0)
int next_count = 0;
```

- n Each procedure *F* will be replaced by

```
wait(mutex);
...
body of F;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex);
```

- n Mutual exclusion within a monitor is ensured





Monitor Implementation – Condition Variables

- n For each condition variable x , we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

- n The operation $x.wait$ can be implemented as:

```
x_count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x_count--;
```



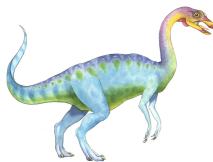


Monitor Implementation (Cont.)

- n The operation `x.signal` can be implemented as:

```
if (x_count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```





Resuming Processes within a Monitor

- n If several processes queued on condition x, and x.signal() executed, which should be resumed?
- n FCFS frequently not adequate
- n **conditional-wait** construct of the form x.wait(c)
 - | Where c is **priority number**
 - | Process with lowest number (highest priority) is scheduled next





Single Resource allocation

- n Allocate a single resource among competing processes using priority numbers that specify the maximum time a process plans to use the resource

```
R.acquire(t);  
...  
access the resource;  
...  
  
R.release;
```

- n Where R is an instance of type **ResourceAllocator**





A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```





Synchronization Examples

- n Solaris
- n Windows
- n Linux
- n Pthreads





Solaris Synchronization

- n Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing
- n Uses **adaptive mutexes** for efficiency when protecting data from short code segments
 - | Starts as a standard semaphore spin-lock
 - | If lock held, and by a thread running on another CPU, spins
 - | If lock held by non-run-state thread, block and sleep waiting for signal of lock being released
- n Uses **condition variables**
- n Uses **readers-writers** locks when longer sections of code need access to data
- n Uses **turnstile**s to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock
 - | Turnstiles are per-lock-holding-thread, not per-object
- n Priority-inheritance per-turnstile gives the running thread the highest of the priorities of the threads in its turnstile





Windows Synchronization

- n Uses interrupt masks to protect access to global resources on uniprocessor systems
- n Uses **spinlocks** on multiprocessor systems
 - | Spinlocking-thread will never be preempted
- n Also provides **dispatcher objects** user-land which may act mutexes, semaphores, events, and timers
 - | **Events**
 - ▶ An event acts much like a condition variable
 - | Timers notify one or more thread when time expired
 - | Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (thread will block)

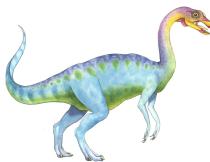




Linux Synchronization

- n Linux:
 - | Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - | Version 2.6 and later, fully preemptive
- n Linux provides:
 - | Semaphores
 - | atomic integers
 - | spinlocks
 - | reader-writer versions of both
- n On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption





Pthreads Synchronization

- n Pthreads API is OS-independent
- n It provides:
 - | mutex locks
 - | condition variable
- n Non-portable extensions include:
 - | read-write locks
 - | spinlocks





Alternative Approaches

- n Transactional Memory
- n OpenMP
- n Functional Programming Languages





Transactional Memory

- A **memory transaction** is a sequence of read-write operations to memory that are performed atomically.

```
void update()
{
    /* read/write memory */
}
```





OpenMP

- OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

The code contained within the `#pragma omp critical` directive is treated as a critical section and performed atomically.





Functional Programming Languages

- n Functional programming languages offer a different paradigm than procedural languages in that they do not maintain state.
- n Variables are treated as immutable and cannot change state once they have been assigned a value.
- n There is increasing interest in functional languages such as Erlang and Scala for their approach in handling data races.



End of Chapter 5

