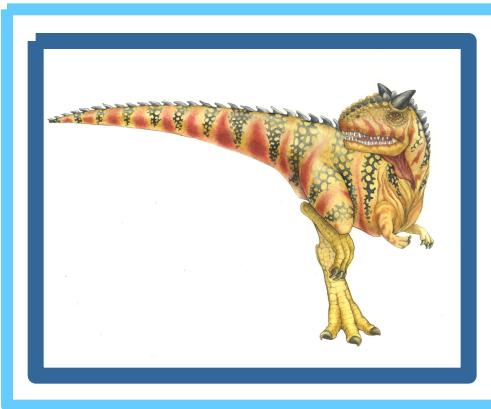


Chapter 8: Main Memory

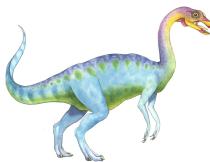




Chapter 8: Memory Management

- n Background
- n Swapping
- n Contiguous Memory Allocation
- n Segmentation
- n Paging
- n Structure of the Page Table
- n Example: The Intel 32 and 64-bit Architectures
- n Example: ARM Architecture

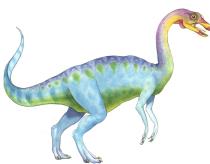




Objectives

- n To provide a detailed description of various ways of organizing memory hardware
- n To discuss various memory-management techniques, including paging and segmentation
- n To provide a detailed description of the Intel Pentium, which supports both pure segmentation and segmentation with paging





Background

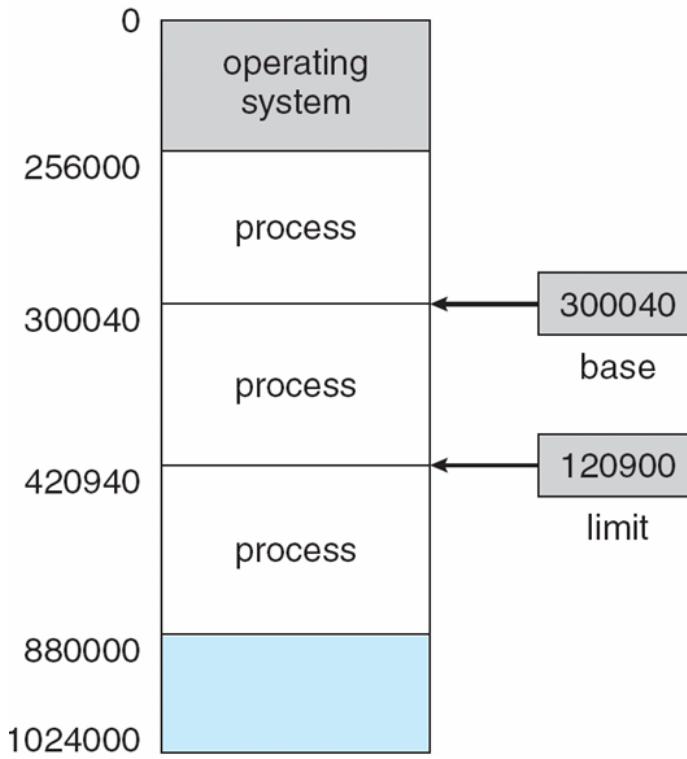
- n Program must be brought (from disk) into memory and placed within a process for it to be run
- n Main memory and registers are only storage CPU can access directly
- n Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- n Register access in one CPU clock (or less)
- n Main memory can take many cycles, causing a **stall**
- n **Cache** sits between main memory and CPU registers
- n Protection of memory required to ensure correct operation





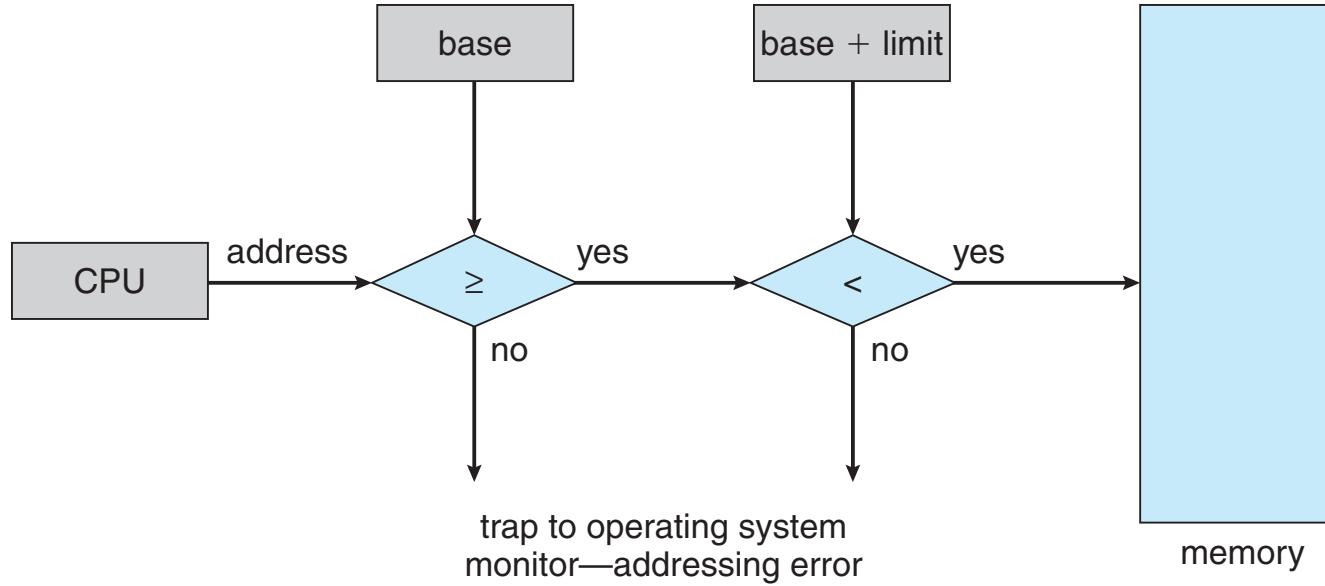
Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user





Hardware Address Protection





Address Binding

- n Programs on disk, ready to be brought into memory to execute from an **input queue**
 - | Without support, must be loaded into address 0000
- n Inconvenient to have first user process physical address always at 0000
 - | How can it not be?
- n Further, addresses represented in different ways at different stages of a program's life
 - | Source code addresses usually symbolic
 - | Compiled code addresses **bind** to relocatable addresses
 - ▶ i.e. "14 bytes from beginning of this module"
 - | Linker or loader will bind relocatable addresses to absolute addresses
 - ▶ i.e. 74014
 - | Each binding maps one address space to another

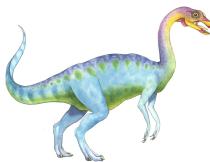




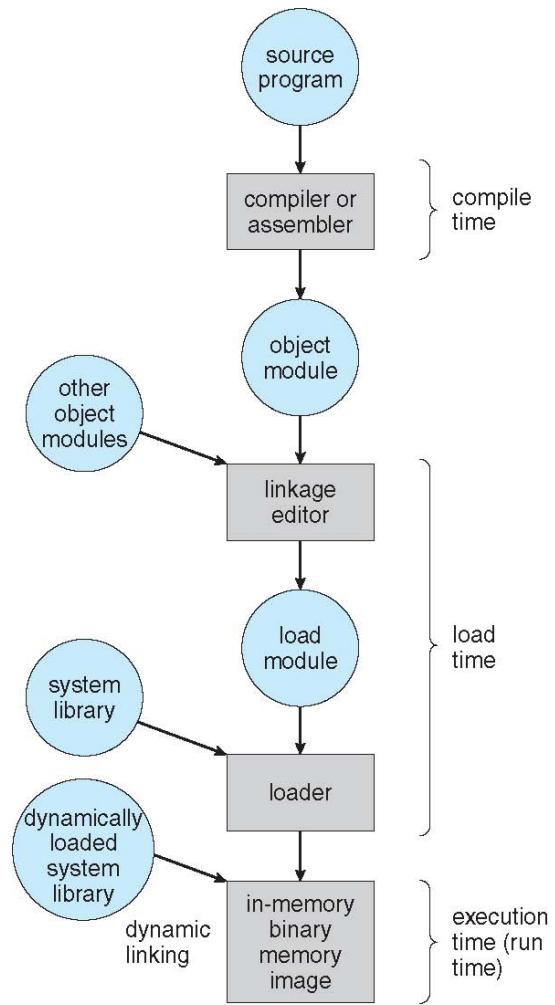
Binding of Instructions and Data to Memory

- n Address binding of instructions and data to memory addresses can happen at three different stages
 - | **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - | **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - | **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ▶ Need hardware support for address maps (e.g., base and limit registers)





Multistep Processing of a User Program





Logical vs. Physical Address Space

- n The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - | **Logical address** – generated by the CPU; also referred to as **virtual address**
 - | **Physical address** – address seen by the memory unit
- n Logical and physical addresses are the same in compile-time and load-time address-binding schemes; logical (virtual) and physical addresses differ in execution-time address-binding scheme
- n **Logical address space** is the set of all logical addresses generated by a program
- n **Physical address space** is the set of all physical addresses generated by a program





Memory-Management Unit (MMU)

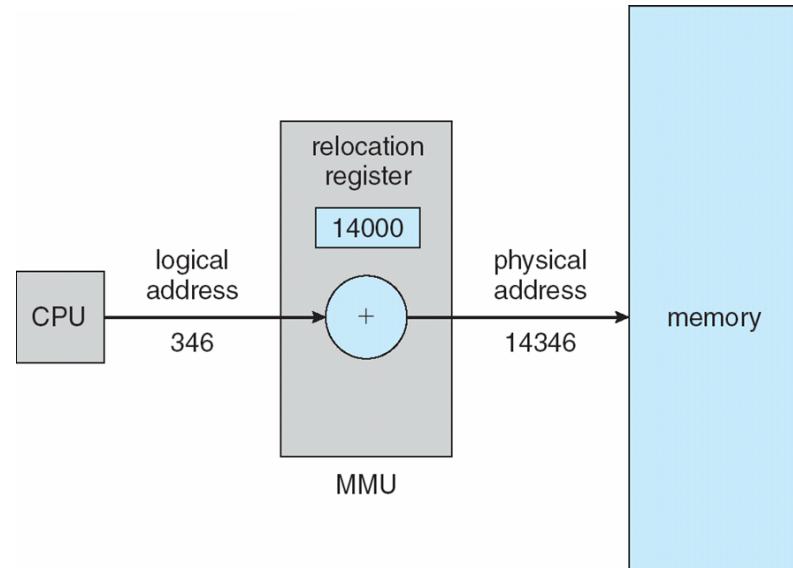
- n Hardware device that at run time maps virtual to physical address
- n Many methods possible, covered in the rest of this chapter
- n To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - | Base register now called **relocation register**
 - | MS-DOS on Intel 80x86 used 4 relocation registers
- n The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - | Execution-time binding occurs when reference is made to location in memory
 - | Logical address bound to physical addresses





Dynamic relocation using a relocation register

- n Routine is not loaded until it is called
- n Better memory-space utilization; unused routine is never loaded
- n All routines kept on disk in relocatable load format
- n Useful when large amounts of code are needed to handle infrequently occurring cases
- n No special support from the operating system is required
 - | Implemented through program design
 - | OS can help by providing libraries to implement dynamic loading





Dynamic Linking

- n **Static linking** – system libraries and program code combined by the loader into the binary program image
- n Dynamic linking –linking postponed until execution time
- n Small piece of code, **stub**, used to locate the appropriate memory-resident library routine
- n Stub replaces itself with the address of the routine, and executes the routine
- n Operating system checks if routine is in processes' memory address
 - | If not in address space, add to address space
- n Dynamic linking is particularly useful for libraries
- n System also known as **shared libraries**
- n Consider applicability to patching system libraries
 - | Versioning may be needed

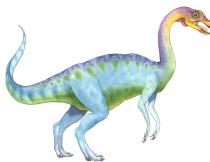




Swapping

- n A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
 - | Total physical memory space of processes can exceed physical memory
- n **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- n **Roll out, roll in** – swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- n Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped
- n System maintains a **ready queue** of ready-to-run processes which have memory images on disk





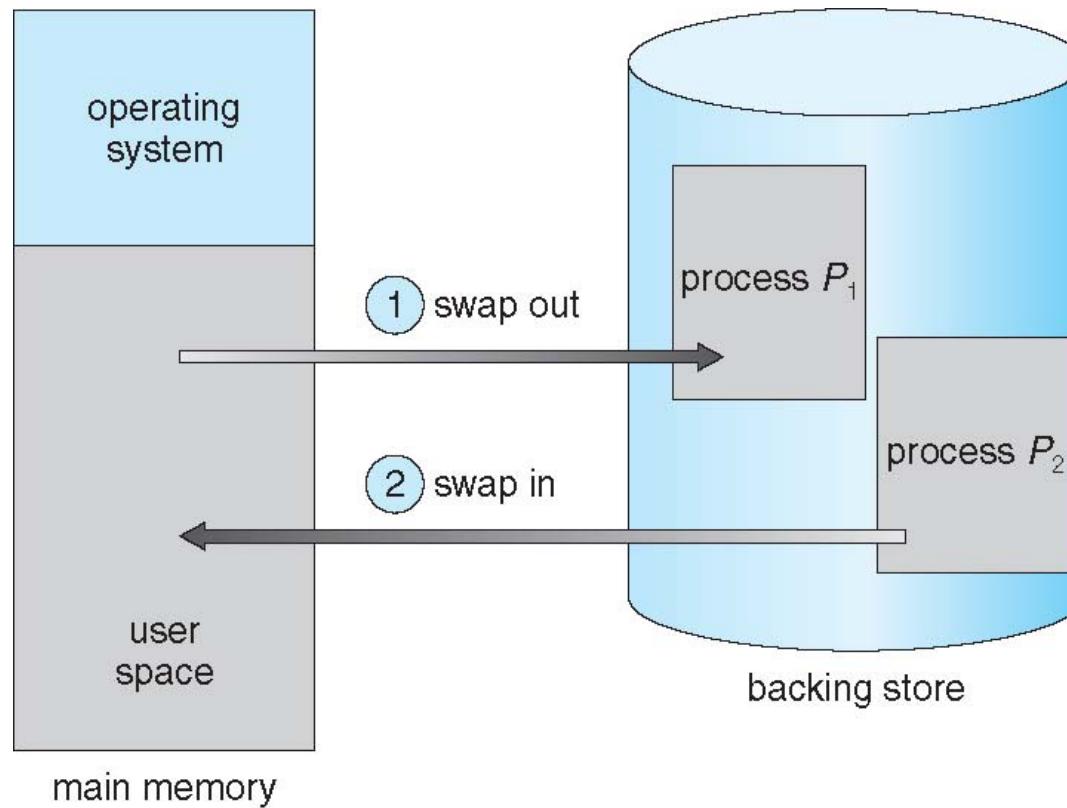
Swapping (Cont.)

- n Does the swapped out process need to swap back in to same physical addresses?
- n Depends on address binding method
 - | Plus consider pending I/O to / from process memory space
- n Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - | Swapping normally disabled
 - | Started if more than threshold amount of memory allocated
 - | Disabled again once memory demand reduced below threshold





Schematic View of Swapping





Context Switch Time including Swapping

- n If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- n Context switch time can then be very high
- n 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - | Swap out time of 2000 ms
 - | Plus swap in of same sized process
 - | Total context switch swapping component time of 4000ms (4 seconds)
- n Can reduce if reduce size of memory swapped – by knowing how much memory really being used
 - | System calls to inform OS of memory use via `request_memory()` and `release_memory()`





Context Switch Time and Swapping (Cont.)

- n Other constraints as well on swapping
 - | Pending I/O – can't swap out as I/O would occur to wrong process
 - | Or always transfer I/O to kernel space, then to I/O device
 - ▶ Known as **double buffering**, adds overhead
- n Standard swapping not used in modern operating systems
 - | But modified version common
 - ▶ Swap only when free memory extremely low





Swapping on Mobile Systems

- n Not typically supported
 - | Flash memory based
 - ▶ Small amount of space
 - ▶ Limited number of write cycles
 - ▶ Poor throughput between flash memory and CPU on mobile platform
- n Instead use other methods to free memory if low
 - | iOS **asks** apps to voluntarily relinquish allocated memory
 - ▶ Read-only data thrown out and reloaded from flash if needed
 - ▶ Failure to free can result in termination
 - | Android terminates apps if low free memory, but first writes **application state** to flash for fast restart
 - | Both OSes support paging as discussed below





Contiguous Allocation

- n Main memory must support both OS and user processes
- n Limited resource, must allocate efficiently
- n Contiguous allocation is one early method
- n Main memory usually into two **partitions**:
 - | Resident operating system, usually held in low memory with interrupt vector
 - | User processes then held in high memory
 - | Each process contained in single contiguous section of memory





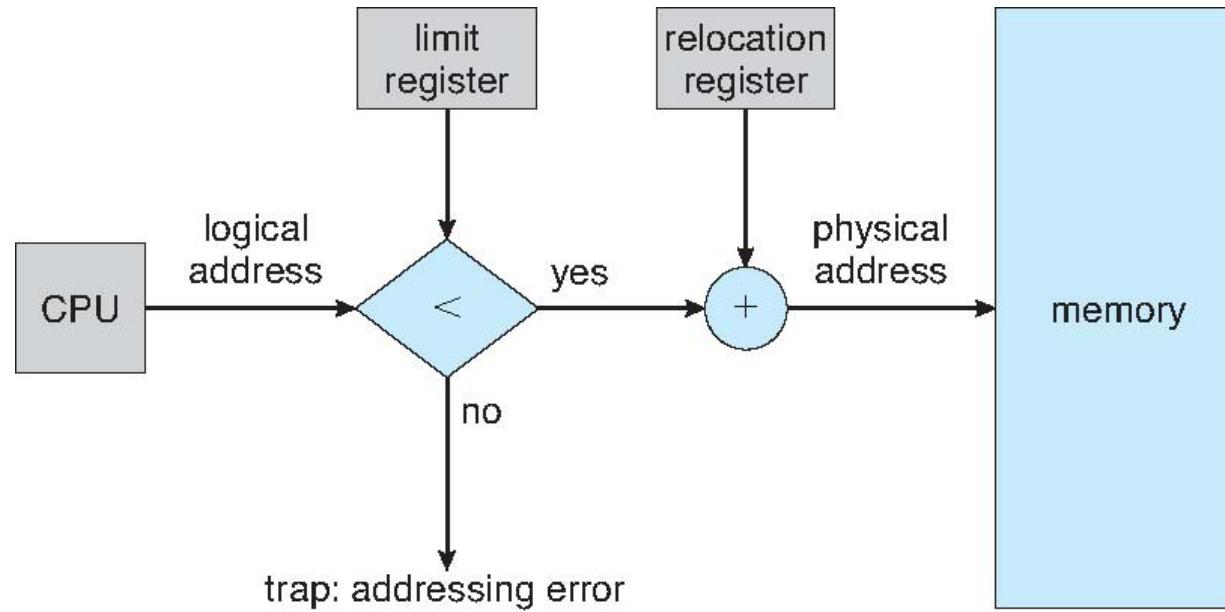
Contiguous Allocation (Cont.)

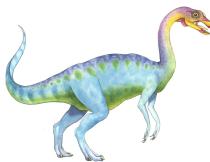
- n Relocation registers used to protect user processes from each other, and from changing operating-system code and data
 - | Base register contains value of smallest physical address
 - | Limit register contains range of logical addresses – each logical address must be less than the limit register
 - | MMU maps logical address *dynamically*
 - | Can then allow actions such as kernel code being **transient** and kernel changing size





Hardware Support for Relocation and Limit Registers

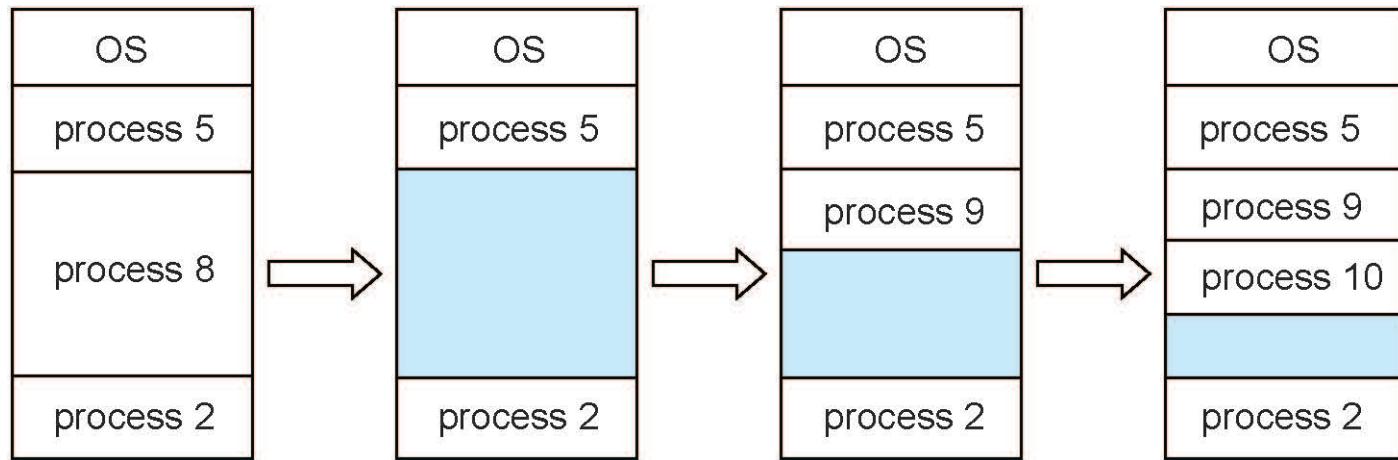




Multiple-partition allocation

n Multiple-partition allocation

- | Degree of multiprogramming limited by number of partitions
- | **Variable-partition** sizes for efficiency (sized to a given process' needs)
- | **Hole** – block of available memory; holes of various size are scattered throughout memory
- | When a process arrives, it is allocated memory from a hole large enough to accommodate it
- | Process exiting frees its partition, adjacent free partitions combined
- | Operating system maintains information about:
 - a) allocated partitions
 - b) free partitions (hole)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

- n **First-fit**: Allocate the *first* hole that is big enough
- n **Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
 - | Produces the smallest leftover hole
- n **Worst-fit**: Allocate the *largest* hole; must also search entire list
 - | Produces the largest leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization





Fragmentation

- n **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- n **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used
- n First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - | 1/3 may be unusable -> **50-percent rule**





Fragmentation (Cont.)

- n Reduce external fragmentation by **compaction**
 - | Shuffle memory contents to place all free memory together in one large block
 - | Compaction is possible *only* if relocation is dynamic, and is done at execution time
 - | I/O problem
 - ▶ Latch job in memory while it is involved in I/O
 - ▶ Do I/O only into OS buffers
- n Now consider that backing store has same fragmentation problems





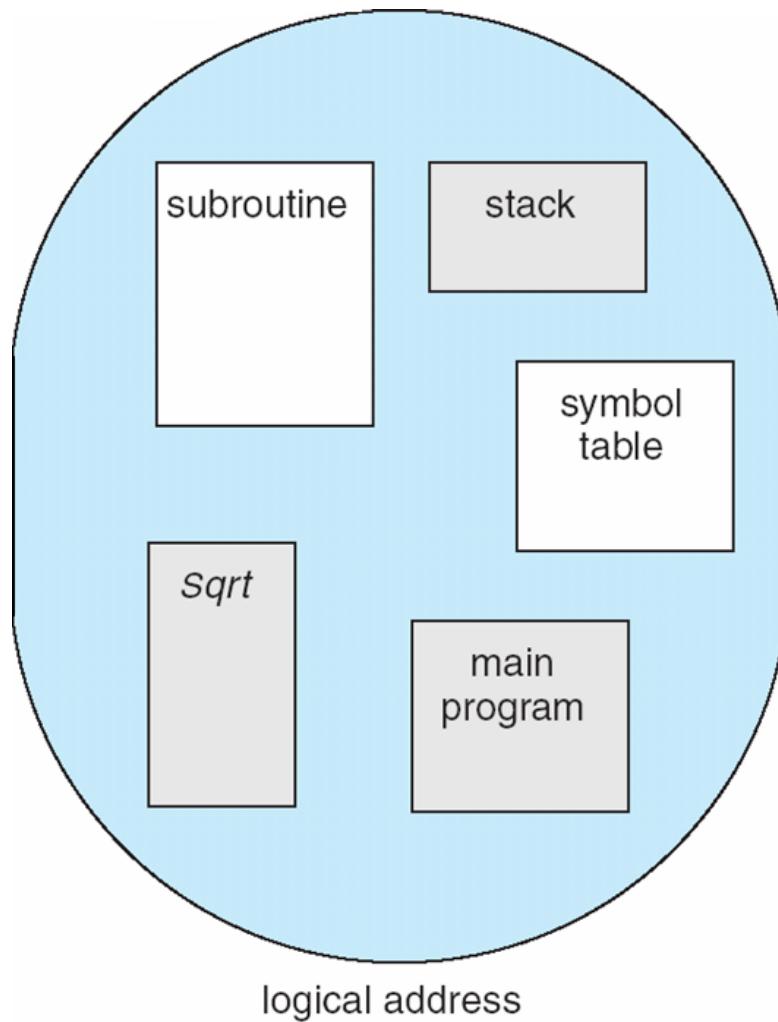
Segmentation

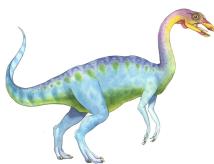
- n Memory-management scheme that supports user view of memory
- n A program is a collection of segments
 - | A segment is a logical unit such as:
 - main program
 - procedure
 - function
 - method
 - object
 - local variables, global variables
 - common block
 - stack
 - symbol table
 - arrays



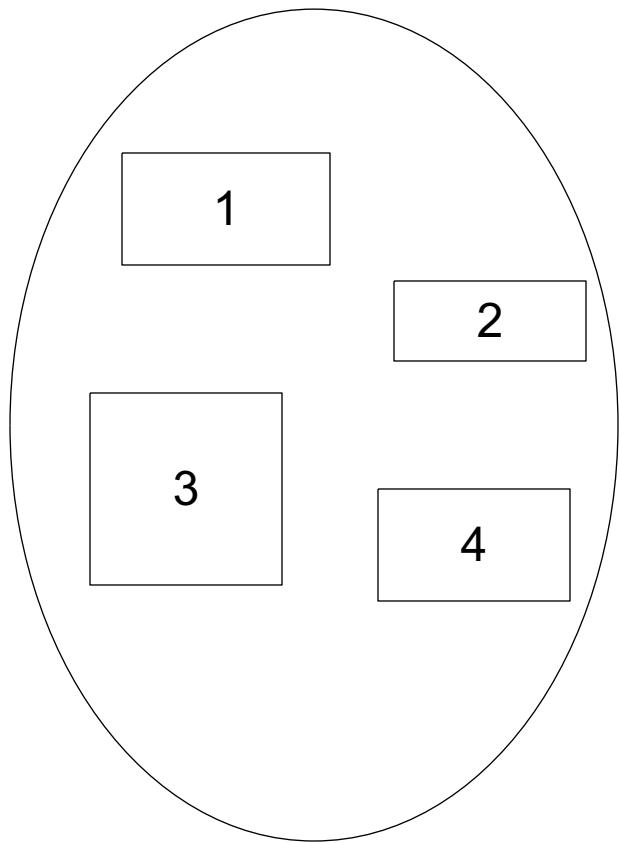


User's View of a Program

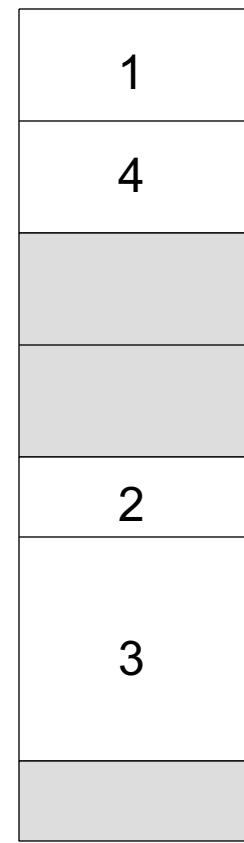




Logical View of Segmentation



user space



physical memory space

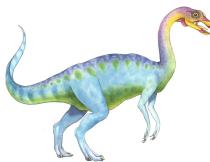




Segmentation Architecture

- n Logical address consists of a two tuple:
 $\langle \text{segment-number}, \text{offset} \rangle$,
- n **Segment table** – maps two-dimensional physical addresses; each table entry has:
 - | **base** – contains the starting physical address where the segments reside in memory
 - | **limit** – specifies the length of the segment
- n **Segment-table base register (STBR)** points to the segment table's location in memory
- n **Segment-table length register (STLR)** indicates number of segments used by a program;
 segment number **s** is legal if **s < STLR**





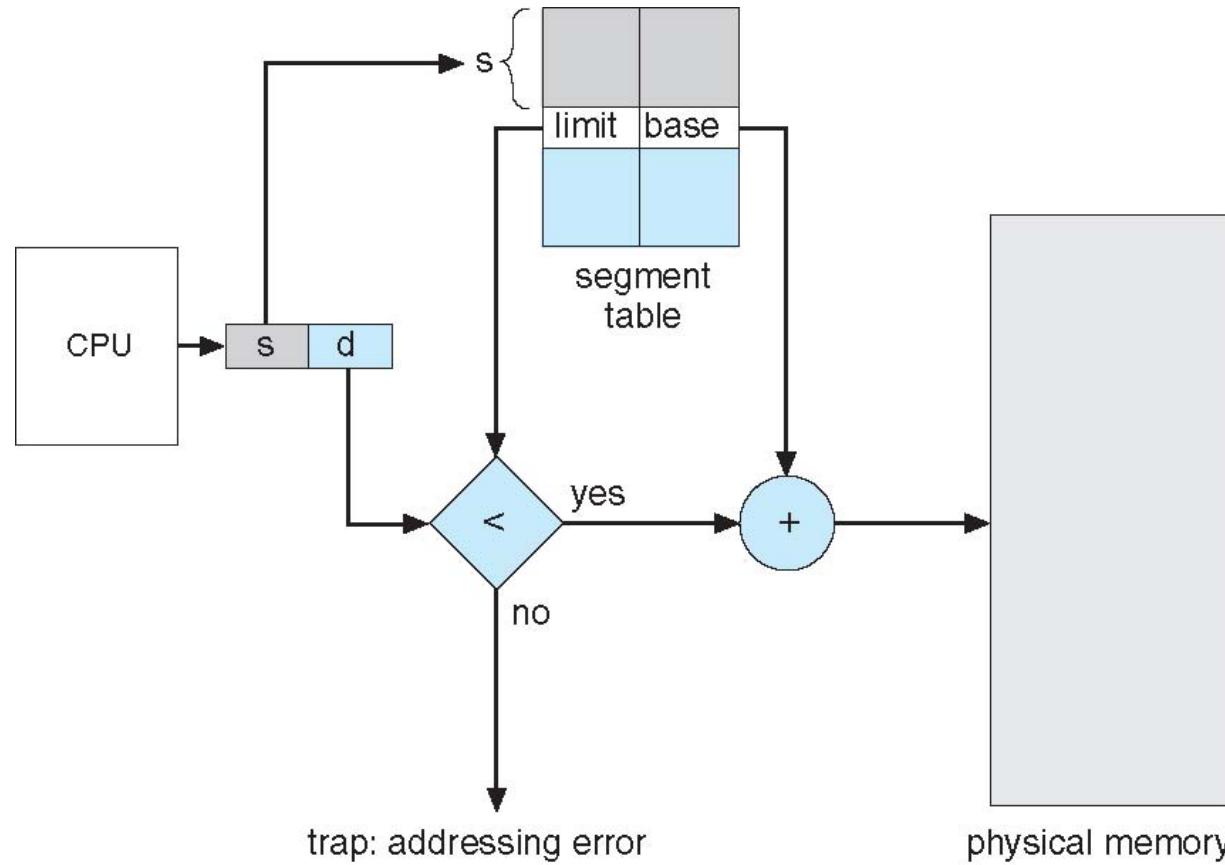
Segmentation Architecture (Cont.)

- n Protection
 - | With each entry in segment table associate:
 - ▶ validation bit = 0 \Rightarrow illegal segment
 - ▶ read/write/execute privileges
- n Protection bits associated with segments; code sharing occurs at segment level
- n Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- n A segmentation example is shown in the following diagram





Segmentation Hardware





Paging

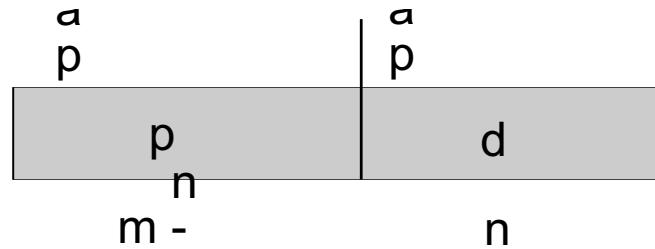
- n Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - | Avoids external fragmentation
 - | Avoids problem of varying sized memory chunks
- n Divide physical memory into fixed-sized blocks called **frames**
 - | Size is power of 2, between 512 bytes and 16 Mbytes
- n Divide logical memory into blocks of same size called **pages**
- n Keep track of all free frames
- n To run a program of size N pages, need to find N free frames and load program
- n Set up a **page table** to translate logical to physical addresses
- n Backing store likewise split into pages
- n Still have Internal fragmentation





Address Translation Scheme

- n Address generated by CPU is divided into:
 - | **Page number (p)** – used as an index into a **page table** which contains base address of each page in physical memory
 - | **Page offset (d)** – combined with base address to define the physical memory address that is sent to the memory unit

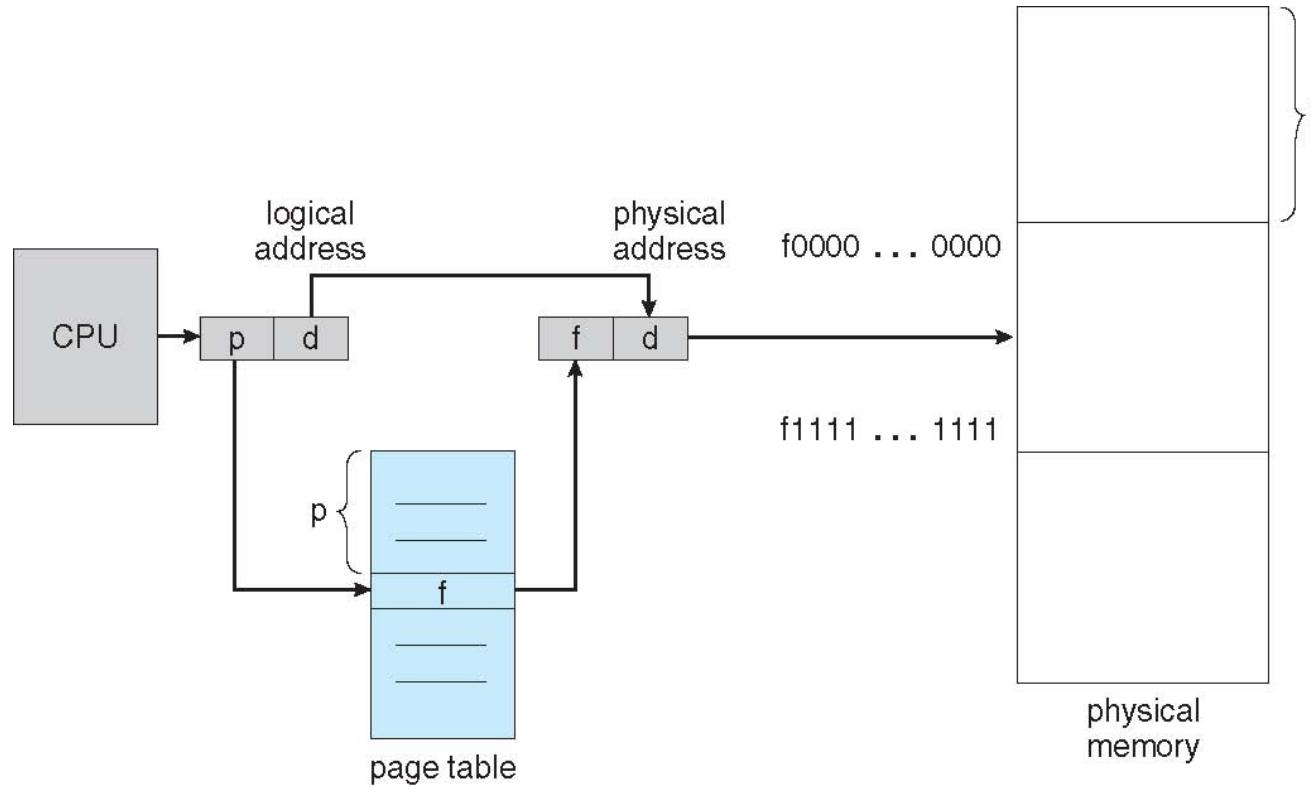


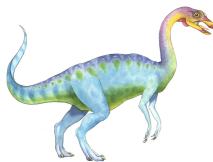
- | For given logical address space 2^m and page size 2^n



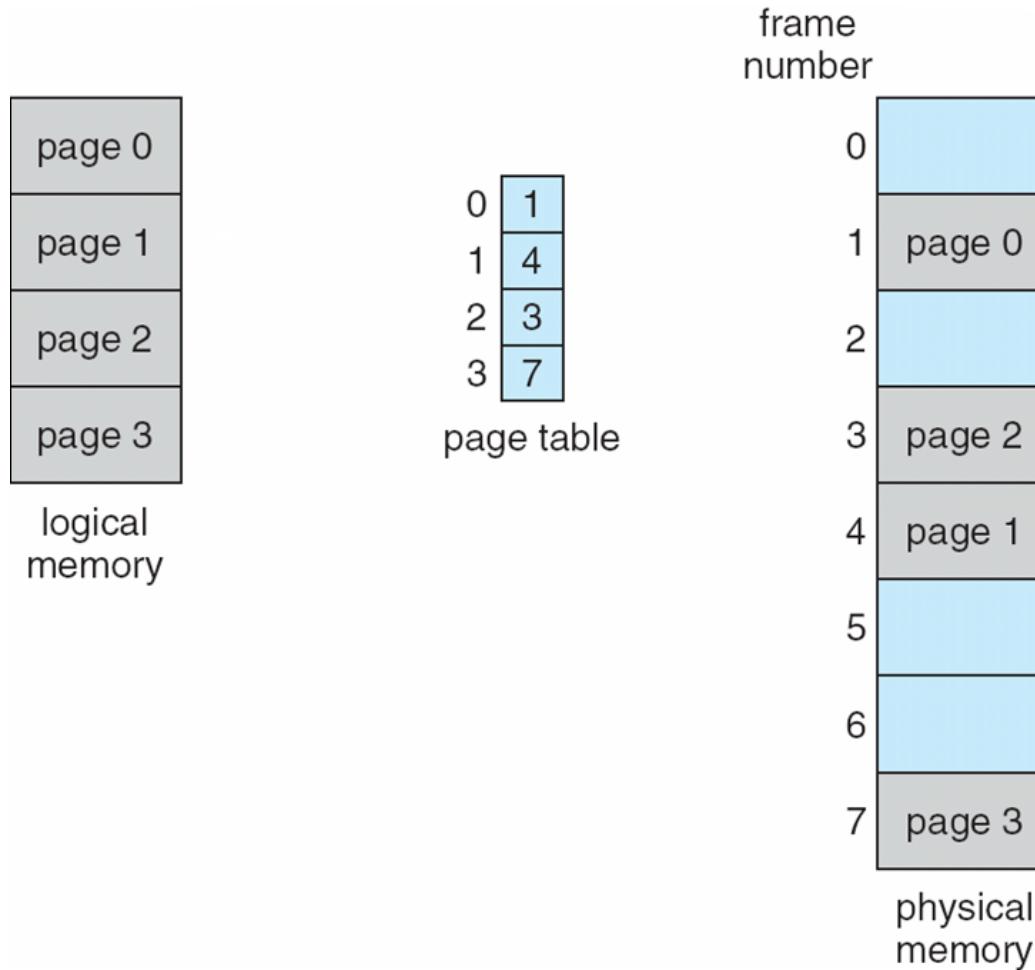


Paging Hardware





Paging Model of Logical and Physical Memory





Paging Example

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table

0	
4	
8	
12	
16	
20	a
21	b
22	c
23	d
24	e
25	f
26	g
27	h
28	

physical memory

$n=2$ and $m=4$ 32-byte memory and 4-byte pages





Paging (Cont.)

- n Calculating internal fragmentation
 - | Page size = 2,048 bytes
 - | Process size = 72,766 bytes
 - | 35 pages + 1,086 bytes
 - | Internal fragmentation of $2,048 - 1,086 = 962$ bytes
 - | Worst case fragmentation = 1 frame – 1 byte
 - | On average fragmentation = $1 / 2$ frame size
 - | So small frame sizes desirable?
 - | But each page table entry takes memory to track
 - | Page sizes growing over time
 - ▶ Solaris supports two page sizes – 8 KB and 4 MB
- n Process view and physical memory now very different
- n By implementation process can only access its own memory

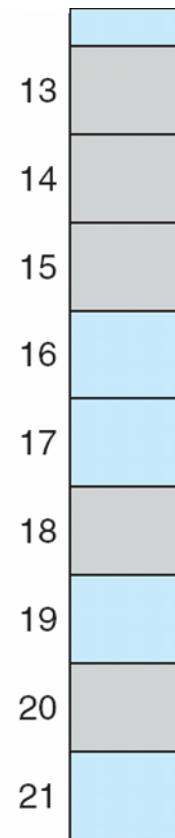
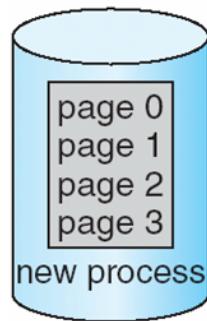




Free Frames

free-frame list

14
13
18
20
15

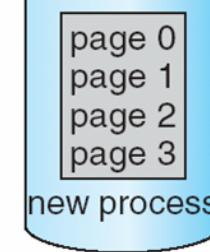


(a)

Before allocation

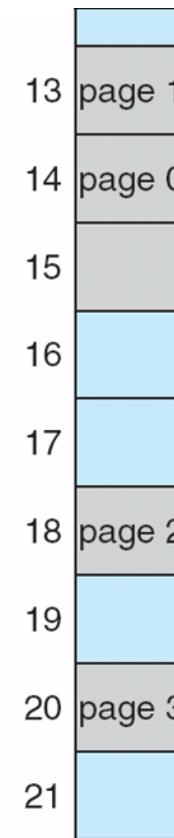
free-frame list

15



0	14
1	13
2	18
3	20

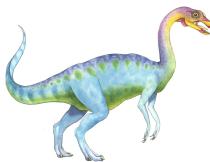
new-process page table



(b)

After allocation





Implementation of Page Table

- n Page table is kept in main memory
- n **Page-table base register (PTBR)** points to the page table
- n **Page-table length register (PTLR)** indicates size of the page table
- n In this scheme every data/instruction access requires two memory accesses
 - | One for the page table and one for the data / instruction
- n The two memory access problem can be solved by the use of a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**





Implementation of Page Table (Cont.)

- n Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - | Otherwise need to flush at every context switch
- n TLBs typically small (64 to 1,024 entries)
- n On a TLB miss, value is loaded into the TLB for faster access next time
 - | Replacement policies must be considered
 - | Some entries can be **wired down** for permanent fast access





Associative Memory

n Associative memory – parallel search

Page #	Frame #

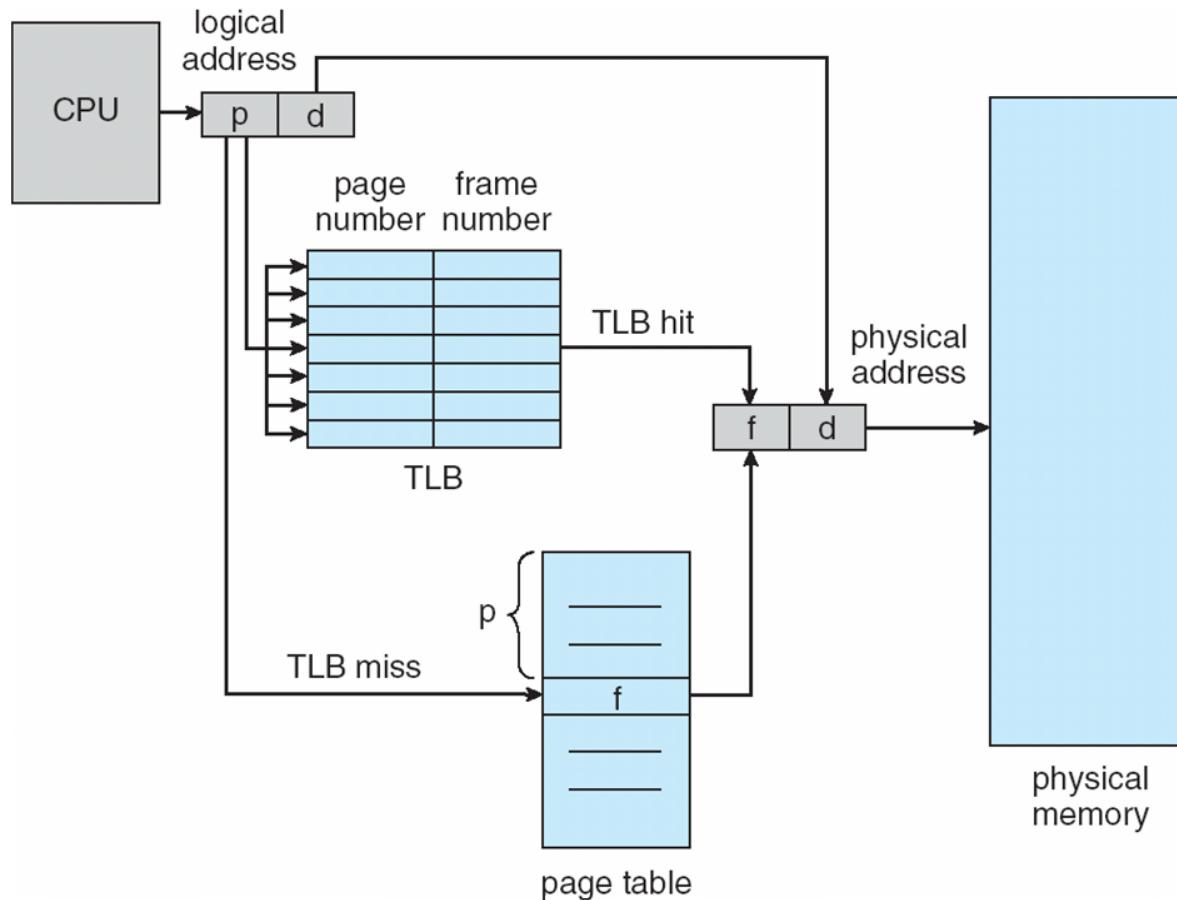
n Address translation (p, d)

- | If p is in associative register, get frame # out
- | Otherwise get frame # from page table in memory





Paging Hardware With TLB

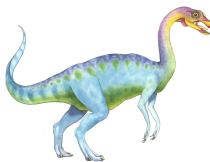




Effective Access Time

- n Associative Lookup = ε time unit
 - | Can be < 10% of memory access time
- n Hit ratio = α
 - | Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers
- n Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
- n **Effective Access Time (EAT)**
$$\begin{aligned} \text{EAT} &= (1 + \varepsilon) \alpha + (2 + \varepsilon)(1 - \alpha) \\ &= 2 + \varepsilon - \alpha \end{aligned}$$
- n Consider $\alpha = 80\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - | $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- n Consider more realistic hit ratio $\rightarrow \alpha = 99\%$, $\varepsilon = 20\text{ns}$ for TLB search, 100ns for memory access
 - | $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$

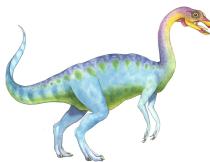




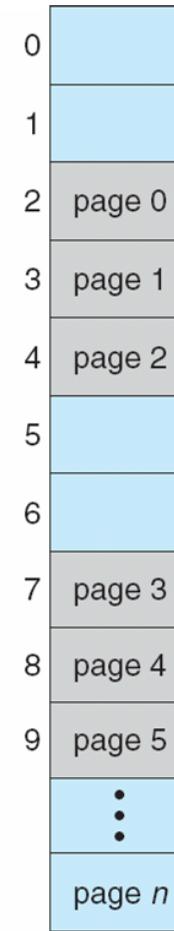
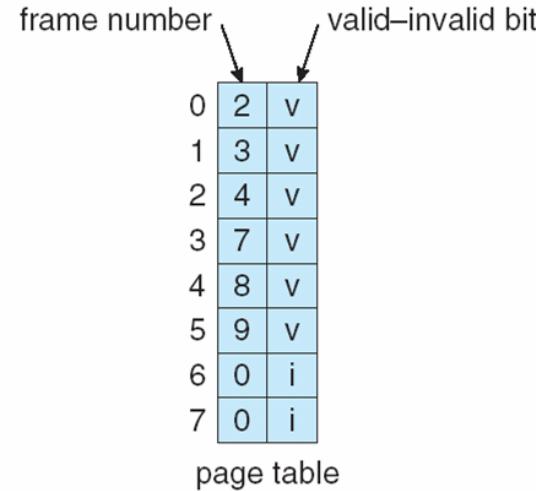
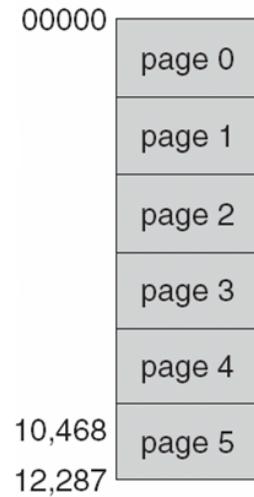
Memory Protection

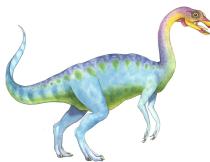
- n Memory protection implemented by associating protection bit with each frame to indicate if read-only or read-write access is allowed
 - | Can also add more bits to indicate page execute-only, and so on
- n **Valid-invalid** bit attached to each entry in the page table:
 - | “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page
 - | “invalid” indicates that the page is not in the process’ logical address space
 - | Or use **page-table length register (PTLR)**
- n Any violations result in a trap to the kernel





Valid (v) or Invalid (i) Bit In A Page Table





Shared Pages

n Shared code

- | One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems)
- | Similar to multiple threads sharing the same process space
- | Also useful for interprocess communication if sharing of read-write pages is allowed

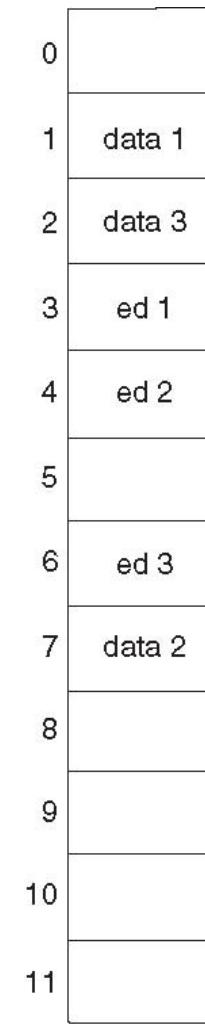
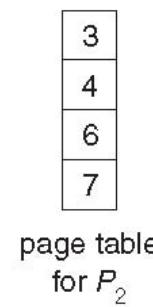
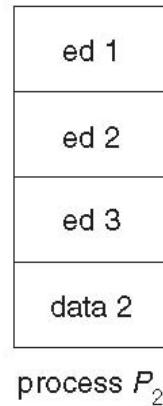
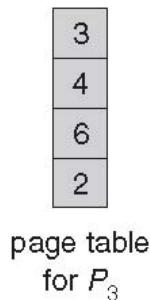
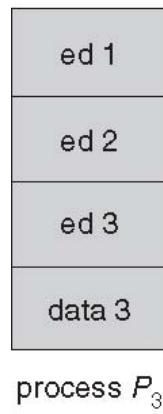
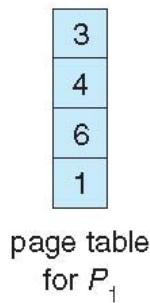
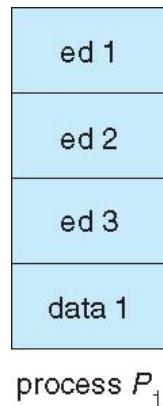
n Private code and data

- | Each process keeps a separate copy of the code and data
- | The pages for the private code and data can appear anywhere in the logical address space





Shared Pages Example

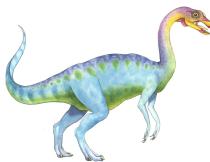




Structure of the Page Table

- n Memory structures for paging can get huge using straight-forward methods
 - | Consider a 32-bit logical address space as on modern computers
 - | Page size of 4 KB (2^{12})
 - | Page table would have 1 million entries ($2^{32} / 2^{12}$)
 - | If each entry is 4 bytes -> 4 MB of physical address space / memory for page table alone
 - ▶ That amount of memory used to cost a lot
 - ▶ Don't want to allocate that contiguously in main memory
- n Hierarchical Paging
- n Hashed Page Tables
- n Inverted Page Tables





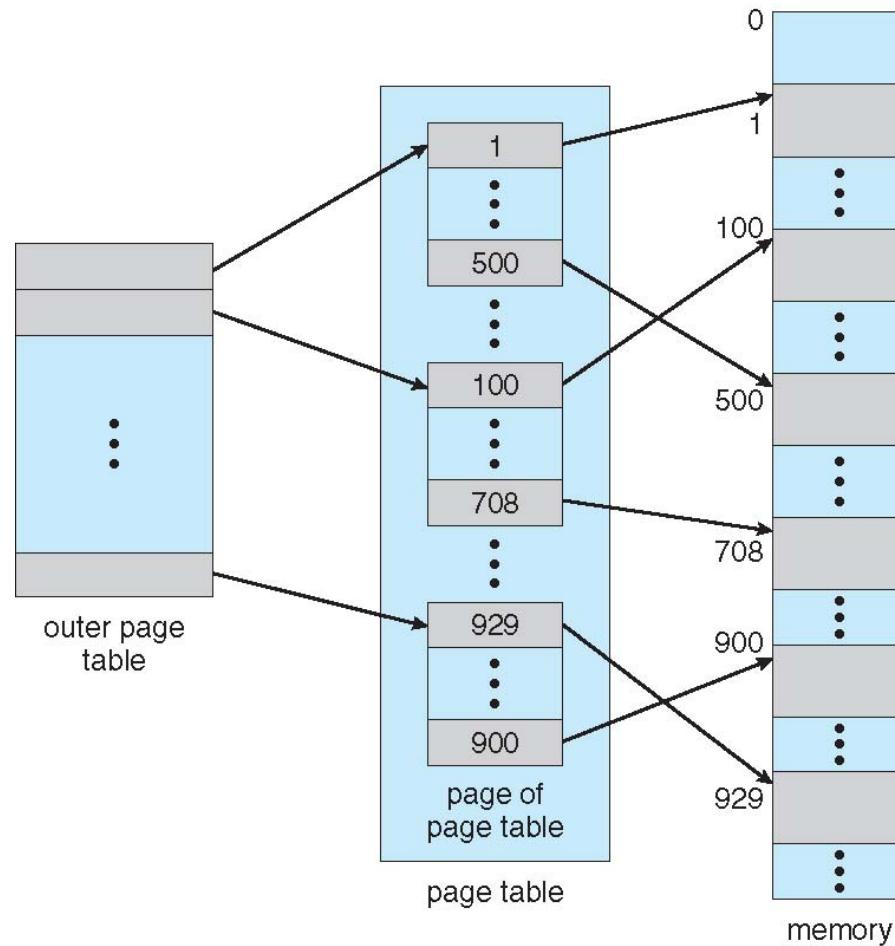
Hierarchical Page Tables

- n Break up the logical address space into multiple page tables
- n A simple technique is a two-level page table
- n We then page the page table





Two-Level Page-Table Scheme





Two-Level Paging Example

- n A logical address (on 32-bit machine with 1K page size) is divided into:
 - | a page number consisting of 22 bits
 - | a page offset consisting of 10 bits
- n Since the page table is paged, the page number is further divided into:
 - | a 12-bit page number
 - | a 10-bit page offset
- n Thus, a logical address is as follows:

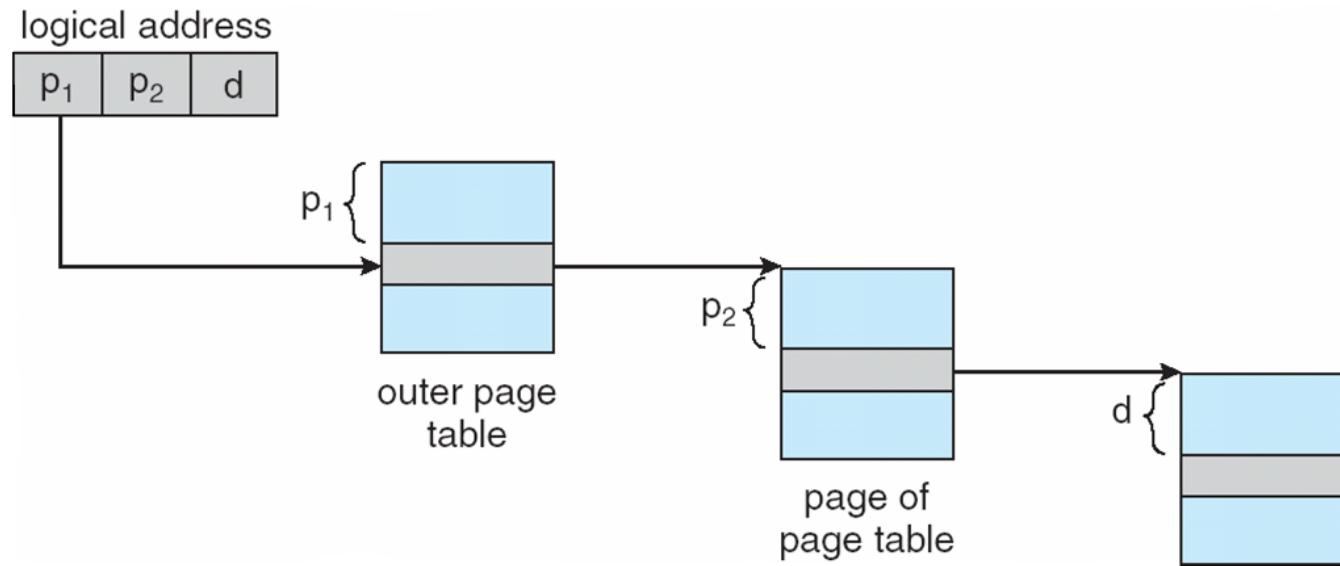
page number	page offset
p_1	p_2
12	10

- n where p_1 is an index into the outer page table, and p_2 is the displacement within the page of the inner page table
- n Known as **forward-mapped page table**





Address-Translation Scheme





64-bit Logical Address Space

- n Even two-level paging scheme not sufficient
- n If page size is 4 KB (2^{12})
 - | Then page table has 2^{52} entries
 - | If two level scheme, inner page tables could be 2^{10} 4-byte entries
 - | Address would look like

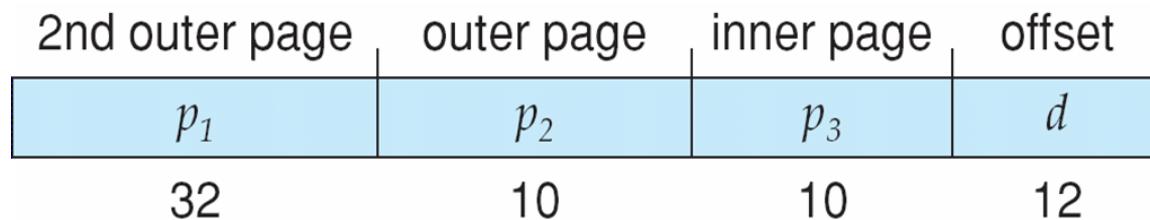
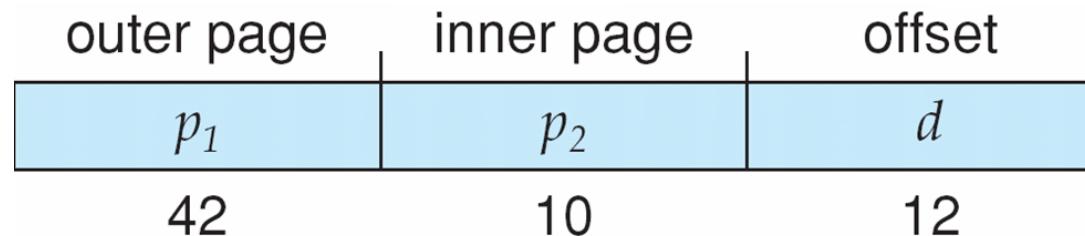
outer page	inner page	page offset
p_1	p_2	d

42 10 12
 - | Outer page table has 2^{42} entries or 2^{44} bytes
 - | One solution is to add a 2nd outer page table
 - | But in the following example the 2nd outer page table is still 2^{34} bytes in size
 - ▶ And possibly 4 memory access to get to one physical memory location





Three-level Paging Scheme





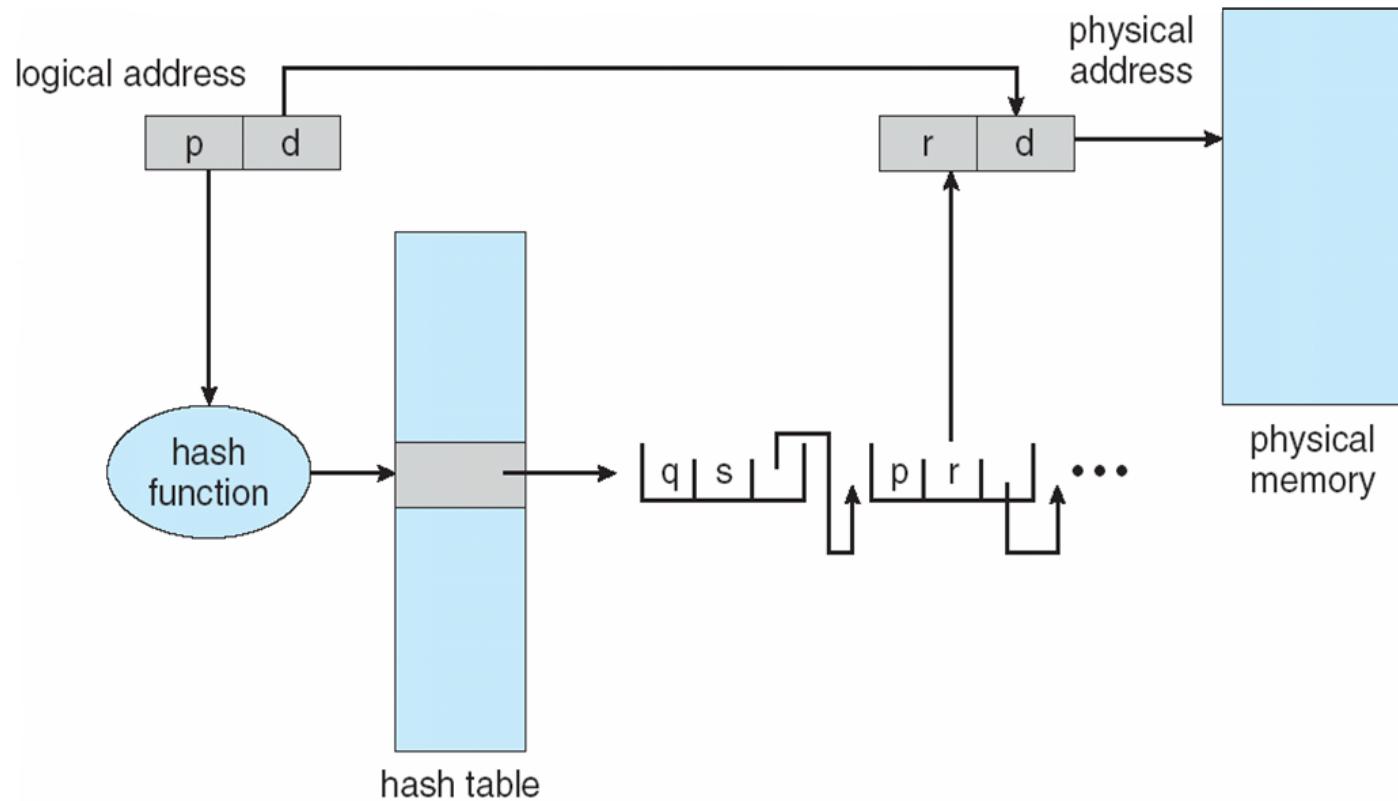
Hashed Page Tables

- n Common in address spaces > 32 bits
- n The virtual page number is hashed into a page table
 - | This page table contains a chain of elements hashing to the same location
- n Each element contains (1) the virtual page number (2) the value of the mapped page frame (3) a pointer to the next element
- n Virtual page numbers are compared in this chain searching for a match
 - | If a match is found, the corresponding physical frame is extracted
- n Variation for 64-bit addresses is **clustered page tables**
 - | Similar to hashed but each entry refers to several pages (such as 16) rather than 1
 - | Especially useful for **sparse** address spaces (where memory references are non-contiguous and scattered)





Hashed Page Table





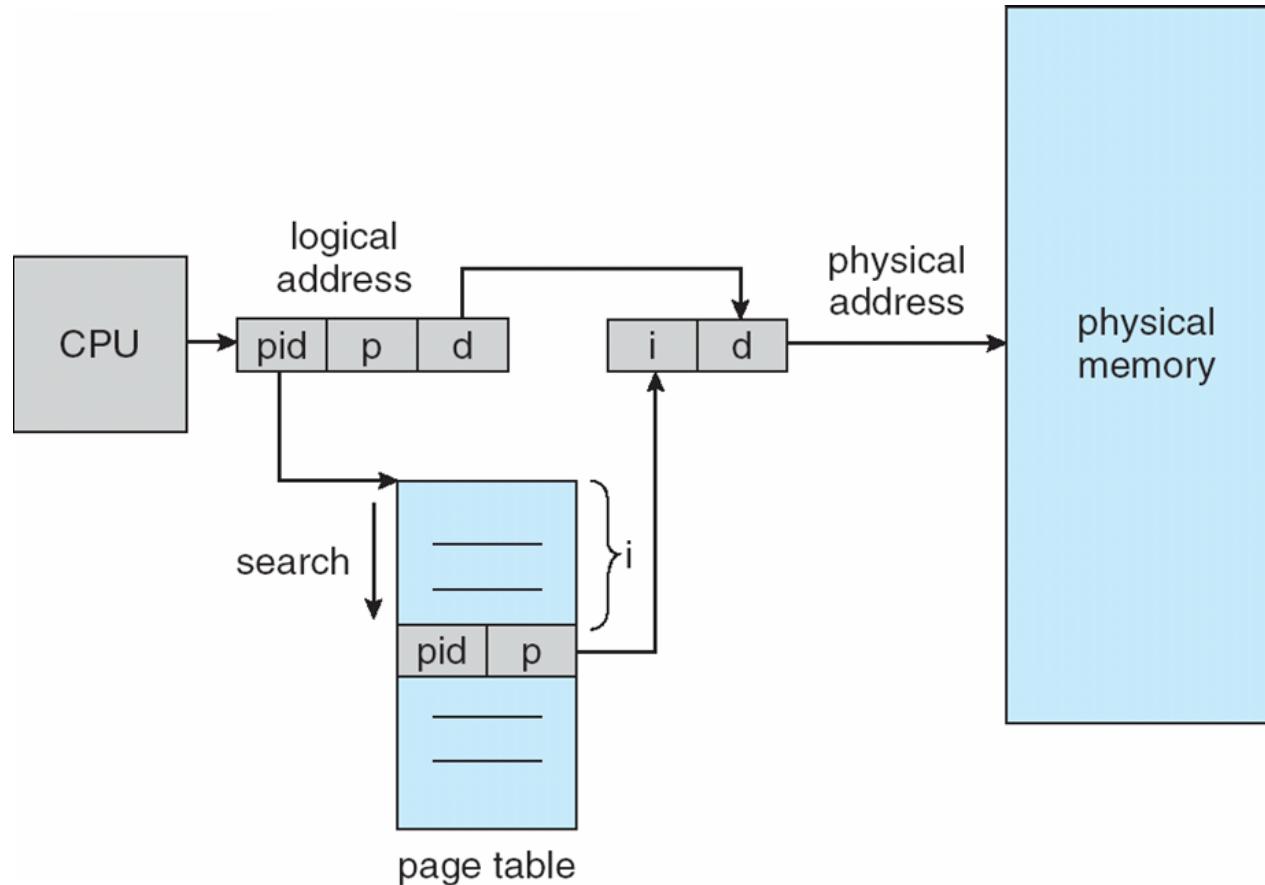
Inverted Page Table

- n Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages
- n One entry for each real page of memory
- n Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page
- n Decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs
- n Use hash table to limit the search to one — or at most a few — page-table entries
 - | TLB can accelerate access
- n But how to implement shared memory?
 - | One mapping of a virtual address to the shared physical address





Inverted Page Table Architecture

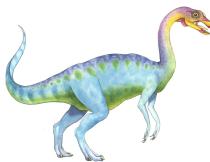




Oracle SPARC Solaris

- n Consider modern, 64-bit operating system example with tightly integrated HW
 - | Goals are efficiency, low overhead
- n Based on hashing, but more complex
- n Two hash tables
 - | One kernel and one for all user processes
 - | Each maps memory addresses from virtual to physical memory
 - | Each entry represents a contiguous area of mapped virtual memory,
 - ▶ More efficient than having a separate hash-table entry for each page
 - | Each entry has base address and span (indicating the number of pages the entry represents)





Oracle SPARC Solaris (Cont.)

- n TLB holds translation table entries (TTEs) for fast hardware lookups
 - | A cache of TTEs reside in a translation storage buffer (TSB)
 - ▶ Includes an entry per recently accessed page
- n Virtual address reference causes TLB search
 - | If miss, hardware walks the in-memory TSB looking for the TTE corresponding to the address
 - ▶ If match found, the CPU copies the TSB entry into the TLB and translation completes
 - ▶ If no match found, kernel interrupted to search the hash table
 - The kernel then creates a TTE from the appropriate hash table and stores it in the TSB, Interrupt handler returns control to the MMU, which completes the address translation.





Example: The Intel 32 and 64-bit Architectures

- n Dominant industry chips
- n Pentium CPUs are 32-bit and called IA-32 architecture
- n Current Intel CPUs are 64-bit and called IA-64 architecture
- n Many variations in the chips, cover the main ideas here





Example: The Intel IA-32 Architecture

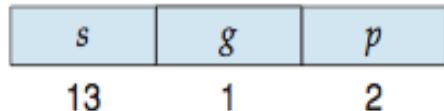
- n Supports both segmentation and segmentation with paging
 - | Each segment can be 4 GB
 - | Up to 16 K segments per process
 - | Divided into two partitions
 - ▶ First partition of up to 8 K segments are private to process (kept in **local descriptor table (LDT)**)
 - ▶ Second partition of up to 8K segments shared among all processes (kept in **global descriptor table (GDT)**)





Example: The Intel IA-32 Architecture (Cont.)

- n CPU generates logical address
 - | Selector given to segmentation unit
 - ▶ Which produces linear addresses

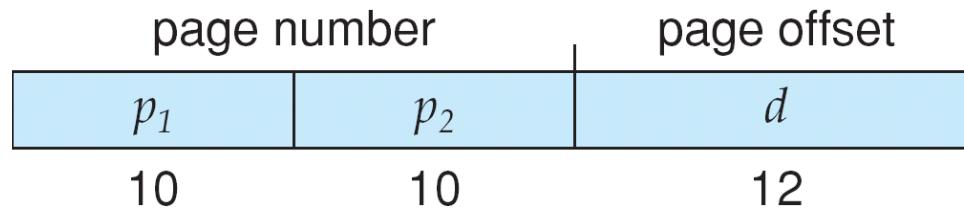
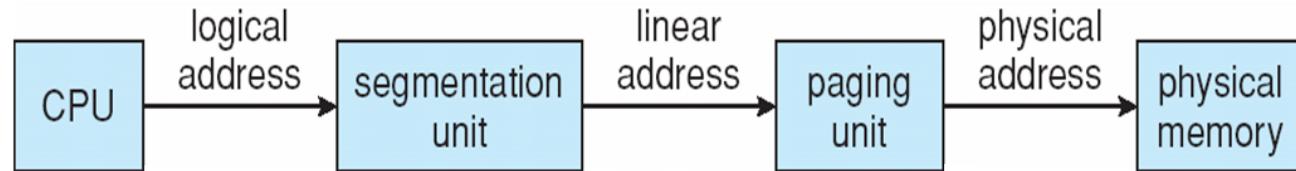


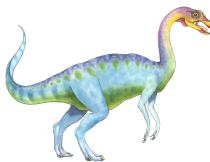
- | Linear address given to paging unit
 - ▶ Which generates physical address in main memory
 - ▶ Paging units form equivalent of MMU
 - ▶ Pages sizes can be 4 KB or 4 MB



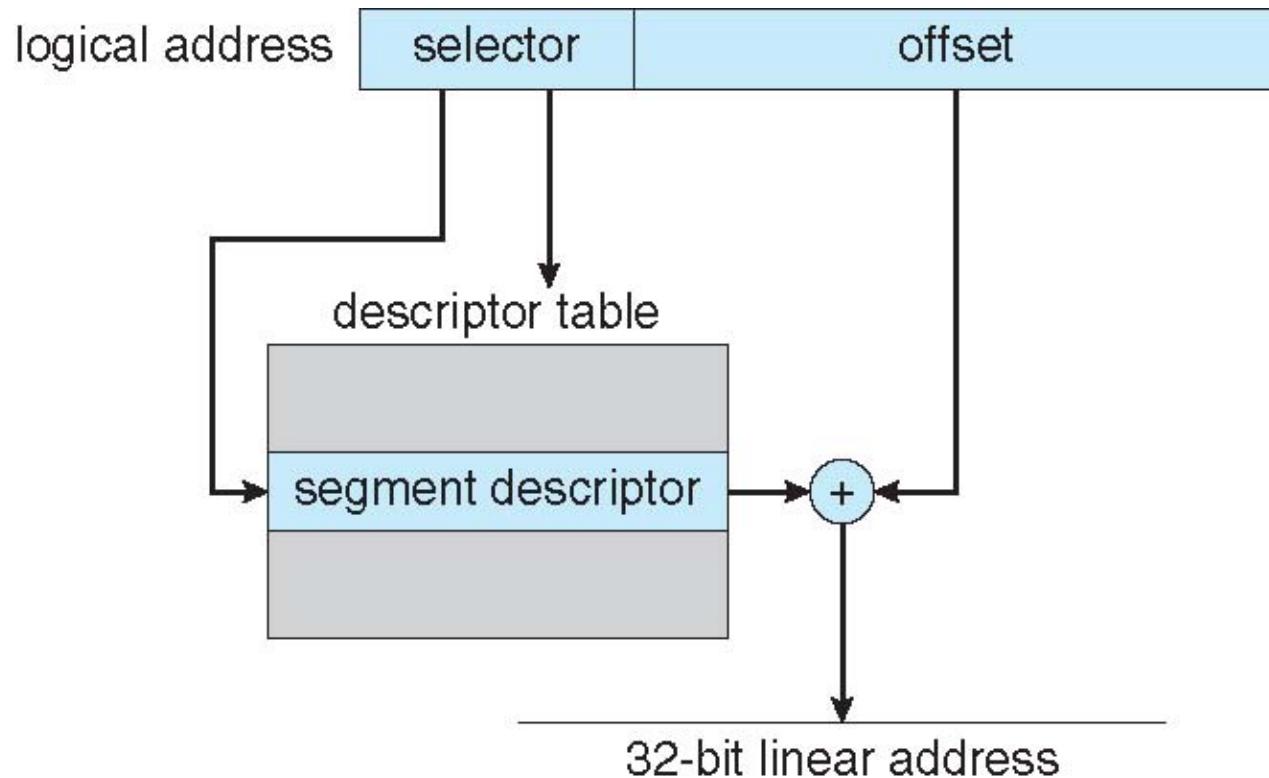


Logical to Physical Address Translation in IA-32



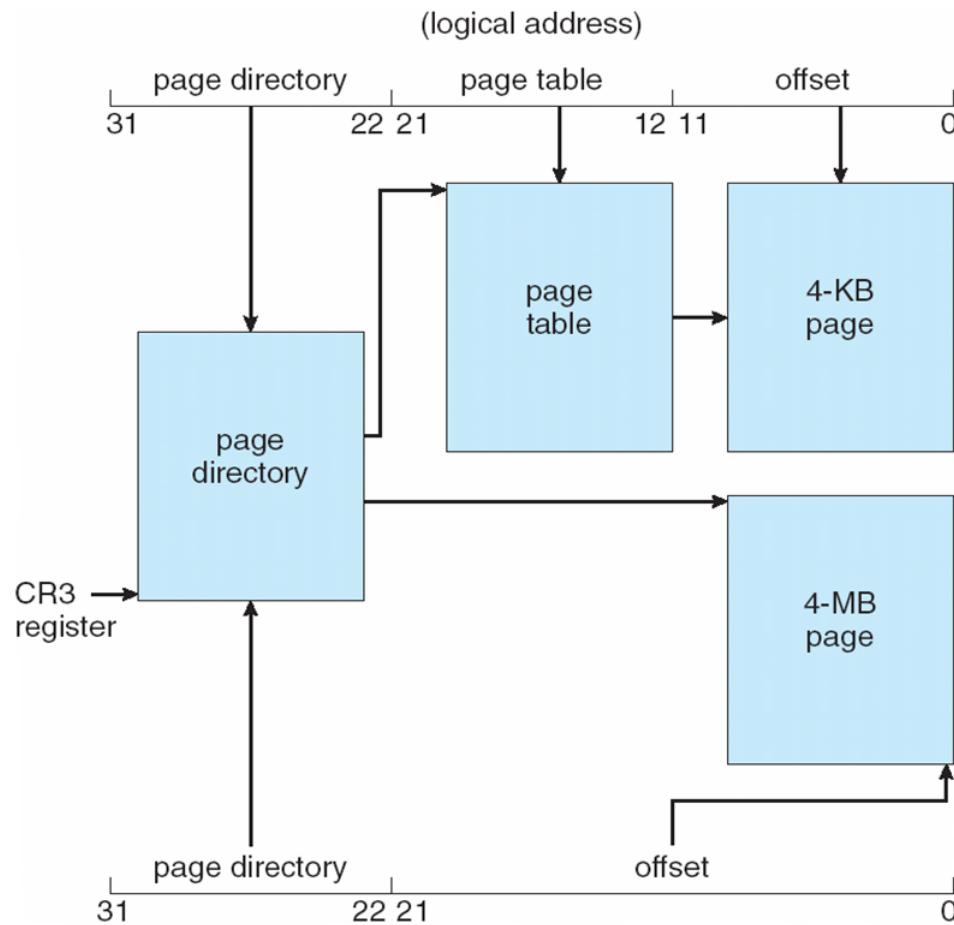


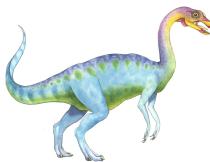
Intel IA-32 Segmentation





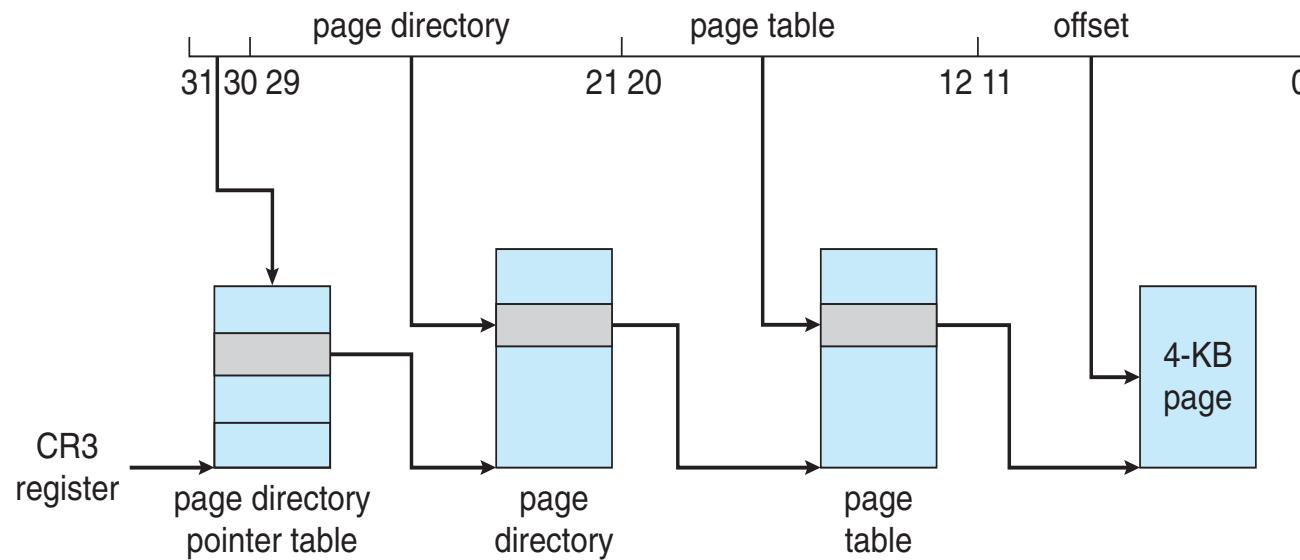
Intel IA-32 Paging Architecture





Intel IA-32 Page Address Extensions

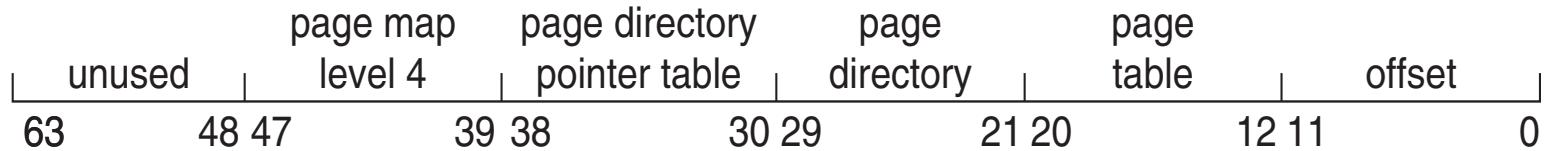
- n 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to more than 4GB of memory space
 - | Paging went to a 3-level scheme
 - | Top two bits refer to a **page directory pointer table**
 - | Page-directory and page-table entries moved to 64-bits in size
 - | Net effect is increasing address space to 36 bits – 64GB of physical memory

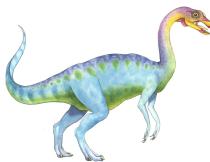




Intel x86-64

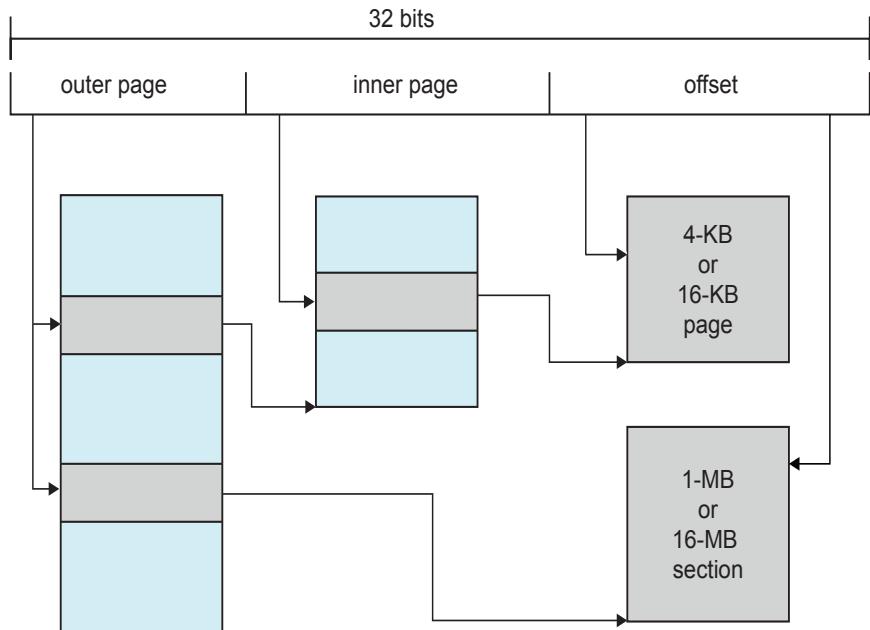
- n Current generation Intel x86 architecture
- n 64 bits is ginormous (> 16 exabytes)
- n In practice only implement 48 bit addressing
 - | Page sizes of 4 KB, 2 MB, 1 GB
 - | Four levels of paging hierarchy
- n Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits





Example: ARM Architecture

- n Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- n Modern, energy efficient, 32-bit CPU
- n 4 KB and 16 KB pages
- n 1 MB and 16 MB pages (termed **sections**)
- n One-level paging for sections, two-level for smaller pages
- n Two levels of TLBs
 - | Outer level has two micro TLBs (one data, one instruction)
 - | Inner is single main TLB
 - | First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



End of Chapter 8

